

Red Hat Ansible Tower MOD

Tzayad Lab Engagement journal

**Red Hat Automation for OCP 3.11 multiple clusters
deployment**

PREPARED FOR: MOD - Eyal David , Guy Aloni

TABLE OF CONTENTS

[TABLE OF CONTENTS](#)

[PREFACE](#)

[CONFIDENTIALITY, COPYRIGHT AND DISCLAIMER](#)

[ABOUT THIS DOCUMENT](#)

[AUDIENCE](#)

[ADDITIONAL BACKGROUND AND RELATED DOCUMENTS](#)

[TERMINOLOGY](#)

[REVISIONS](#)

[PREPARATION](#)

[Staffing](#)

[Engagement Approach](#)

[ENGAGEMENT DETAILS](#)

[Prerequisite Verification](#)

[SOLUTION ARCHITECTURE](#)

[Architecture](#)

[Hardware specifications](#)

[Hardware Detail](#)

[Network Details](#)

[Storage Details](#)

[Supporting Infrastructure](#)

[Supporting Infrastructure](#)

[ENGAGEMENT CLOSURE AND RECOMMENDATIONS.](#)

[Knowledge Transfer](#)

[Testing](#)

[Engagement Observations](#)

[Training and Certification](#)

[Red Hat Certified System Administrator \(RHCSA\)](#)

[RHCSA Rapid Track course with exam \(RH200\)](#)

[Red Hat OpenStack Administration](#)

[APPENDICES](#)

[Appendix A: RHOS Accelerator Test Matrix and Test result](#)

[Appendix B: Installing and configuring Red Hat OpenStack Platform 7.](#)

[Appendix C: Post-deployment configuration.](#)

I. PREFACE

CONFIDENTIALITY, COPYRIGHT AND DISCLAIMER

This is a Customer-facing document between Red Hat, Inc. and MOD:Mamram ("Client").

Copyright © 2015 Red Hat, Inc. All Rights Reserved. No part of the work covered by the copyright herein may be reproduced or used in any form or by any means – graphic, electronic, or mechanical, including photocopying, recording, taping, or information storage and retrieval systems – without permission in writing from Red Hat except as is required to share this information as provided with the aforementioned confidential parties.

This document is not a quote and does not include any binding commitments by Red Hat.

ABOUT THIS DOCUMENT

This is a confidential document between Red Hat, Inc. and MOD: Mamram detailing a Red Hat OCP 3.11 Deployment process Automation by usage of Red Hat Ansible Tower in Disconnected environment and Integration with Central Monitoring/Metrics system (Prometheus/Grafana)

AUDIENCE

The intended audience for this document is any person involved in the review, decision making, sign off or implementation of the proposed solution and Red Hat Consulting engagement to deliver.

ADDITIONAL BACKGROUND AND RELATED DOCUMENTS

This document also references additional information that can be found on Red Hat's documentation site at <https://access.redhat.com/knowledge/docs/> and specifically at

https://access.redhat.com/documentation/en-us/openshift_container_platform/3.11/. &

<https://docs.ansible.com/ansible-tower/3.4.2/html/installandreference/index.html> &

<https://docs.ansible.com/>

Documents specific to products covered in this solution include the following.

- Getting Started Guide
- Installation and Configuration Guide

The following Reference Architecture as in-depth use case studies that may be useful and are available at <https://access.redhat.com/site/articles/reference-architecture> :

- Deploying and Managing OpenShift 3.9 on VMware vSphere
- Automate Red Hat OpenShift Container Platform Deployment on HPE ProLiant Servers with Ansible

TERMINOLOGY

The table below provides a glossary of the terms and acronyms used within this document.

| Term | Definition |
|----------------------------|--|
| RHEL | Red Hat Enterprise Linux. |
| vSphere | VMware Enterprise Virtualization Management appliance |
| ESXi | VMware Enterprise Virtualization Host OS |
| Ansible | Red Hat Automation tool - agentless system configuration tool used for different products deployment and OCP among them |
| Ansible Tower | Red Hat Ansible orchestration product |
| OCP | Red Hat OpenShift Container Platform |
| Disconnected environment | Environment without access to the internet - require preparation of Offline Repository and Offline Container Registry |
| Offline Repository | Web Server in disconnected environment for RPMs repository hosting |
| Offline Container Registry | Asset server with Docker distribution installed in disconnected environment to host container images for OpenStack deployment |
| Inventory | A collection of hosts against which Jobs may be launched. |
| Host | A system managed by Tower, which may include a physical, virtual, cloud-based server, or other device. Typically an operating system instance. Hosts are contained in Inventory. Sometimes referred to as a "node". |
| Group | A set of hosts in Ansible that can be addressed as a set, of which many may exist within a single Inventory. |
| Job Template | The combination of an Ansible playbook and the set of parameters required to launch it. |
| Workflow Job Template | A set consisting of any combination of job templates, project syncs, and inventory syncs, linked together in order to execute them as a single unit. |
| Ad Hoc | Refers to running Ansible to perform some quick command, using <code>/usr/bin/ansible</code> , rather than the orchestration language, which is <code>/usr/bin/ansible-playbook</code> . An example of an ad hoc command might be rebooting 50 machines in your infrastructure. Anything you can do ad hoc can be accomplished by writing a Playbook, and Playbooks can also glue lots of other operations together. |
| Plays | A playbook is a list of plays. A play is minimally a mapping between a set of hosts selected by a host specifier (usually chosen by groups, but sometimes by hostname globs) and the tasks which run on those hosts to define the role that those systems will perform. There can be one or many plays in a playbook. |
| Playbook | An Ansible playbook. Refer to http://docs.ansible.com/ for more information. |
| YAML | Ansible and Tower use YAML to define playbook configuration languages and also variable files. YAML has a minimum of syntax, is very clean, and is easy for people to skim. It is a good data format for configuration files and humans, but is also machine readable. YAML is fairly popular in the dynamic language community and the format has libraries available for serialization in many languages (Python, Perl, Ruby, etc.). |
| Project | A logical collection of Ansible playbooks, represented in Tower. |
| Sudo | Ansible does not require root logins and, since it is daemonless, does not require root level daemons (which can be a security concern in sensitive environments). Ansible can log in and perform many operations wrapped in a sudo command, and can work with both password-less and password-based |

| | |
|--|---|
| | sudo. Some operations that do not normally work with sudo (like scp file transfer) can be achieved with Ansible's copy, template, and fetch modules while running in sudo mode. |
|--|---|

Table I-a: Terminology

REVISIONS

| Name | Version | Date | Comments |
|----------------------------------|---------|------------|------------------------------|
| Anoel Yakoubov, Matan Carmeli | 1.0 | 24.04.2019 | Initial release to customer. |
| | | | |

Table I-b: Revisions

1. PREPARATION

This section details prerequisites and preparation efforts that Client was to complete prior to the engagement. Where applicable, the onsite consultant has validated the provided data and made for any discrepancies or additional detail.

[POC Requirements defined by MOD Tzayad](#)

1.1 Staffing

The following persons were identified by the Client to support the engagement:

| Role | Purpose | Client Assignment | Contact Information |
|------------|--------------------|-------------------|--------------------------|
| Architect | Architecture & DOD | Eyal David | eyaldavid86@gmail.com |
| Consultant | Project Delivery | Matan Carmeli | matan.carmeli7@gmail.com |
| | | | |

Table 1-1: Client staff/contact information.

| Role | Purpose | Red Hat Assignment | Contact Information |
|---------------------|----------------------------|--------------------|---------------------|
| Project Manager | Project Management | Etay Nir | etnir@redhat.com |
| Sr. Cloud Architect | Architecture & Supervision | Orgad Kimchi | okimchi@redhat.com |
| Consultant | Project Delivery | Anoel Yakoubov | anoel@redhat.com |
| Consultant | Project Delivery | Guy Rakover | grakover@redhat.com |

Table 1-2: Red Hat staff/contact information.

1.2 Engagement Approach

This section details Red Hat's strategy to assist MOD: Tzayad with meeting their goals for this engagement. The following table outlines the focus of each phase.

| Purpose | Details |
|--|--|
| VMware vSphere Virtualization environment deployment | <ul style="list-style-type: none"> • Install and Configure ESXi 6.7 as a basis to hosted engine vSphere infrastructure on two physical server • Configure RAID 10 and VMFS as storage backend for vSphere • Install and Configure hosted engine VM appliance on top of ESXi by deploying Windows 2016 VM and from this VM vSphere appliance • Configure vSphere (Data Center, Cluster, Storage and Network Infrastructure) • Install and configure all relevant VMs with RHEL 7.6 for common usage |
| Offline Repository and Offline Container Registry deployment | <ul style="list-style-type: none"> • Offline Repository Installation and configuration on one of VMs to enable Disconnected OCP 3.11.98 deployment • Offline Container Registry Installation and configuration on one of VMs to enable Disconnected OCP 3.11.98 deployment • Configure DNS on this server for OCP environments |
| Ansible Tower Deployment | <ul style="list-style-type: none"> • Deployment and configuration of Red Hat Ansible Tower on one of VMs |
| Central Grafana Server Deployment and Configuration | <ul style="list-style-type: none"> • Installing Grafana on one of VMs. • Configuring Grafana with Data Source - Prometheus |
| Ansible Tower Configurations | <ul style="list-style-type: none"> • Projects, Inventories, Groups, Hosts definitions • Connectivity to Github where all relevant Ansible playbooks and roles resides. • Creation of Job templates and Workflow templates |
| Ansible Tower Templates | <ul style="list-style-type: none"> • Deploy VMware - Job template to create all relevant VMs for OCP deployment • Destroy VMware - Job template to destroy all relevant VMs if condition will be raised • Pre Install - Job template to install different utilities on each VM according preinstall.yml playbook • OCP prerequisites - Job template to run all OCP prerequisites by using openshift-ansible prerequisites.yml playbook • OCP Deploy Cluster - Job template to run all OCP deployment by using openshift-ansible deploy_cluster.yml playbook • Post Jobs - Job template to run all OCP post-deployment configurations, such Grafana deployment etc.. • OCP Deploy Service Catalog - Job template to deploy Service Catalog • VMware and OCP full deployment - Workflow template that orchestrates all Job templates in certain order and conditioning |

Table 1-3: Focus per phase.

2. ENGAGEMENT DETAILS

The engagement focused on the installation and configuration of one (1) environment running VMware Virtualization Platform, on top of which completed Deployment of Ansible Tower Platform and Automation of at least two OCP 3.11 environments Deployment, each with its LB in Disconnected environment. In addition, VMDK used as persistent storage backend for each OCP cluster and each cluster was with its own Metrics system (Prometheus/Grafana) and propagation of metrics to the Central Metrics environment. Optional: GlusterFS usage as persistent storage backend for at least one OCP cluster.

Prerequisite Verification and validation

Prior to commencing the configuration of the lab environments, the following pre-requisites items were validated by our team.

| Criteria | Validated |
|--|-----------|
| Systems are in the Red Hat Hardware Compatibility List for RHEL 7.6. | Validated |
| RHEL 7.6 Server for x86_64 entitlements are available for all systems | Validated |
| VMware entitlements are available for all systems. | Validated |
| OCP 3.11 entitlements are available for all systems. | Validated |
| Compute nodes have enough CPU cores for the expected workload with the required overcommit ratio | Validated |
| Physical nodes for ESXi have CPU with VT-x or AMD-V | Validated |
| Available RAM corresponds to the expected workload and overcommit ratio | Validated |
| Nodes have at least two 1Gb network interfaces | Validated |
| Network topology satisfies requirements for the selected implementation of Ansible Tower and OCP Networking Services | Validated |
| Physical nodes for ESXi have at least 1TB local disk space | Validated |
| RAID controller memory is flash-backed or battery-backed | Validated |

Table 2-1: VMware virtualization nodes hardware verification.

3. SOLUTION ARCHITECTURE

The engagement focused on the installation and configuration of one (1) environment running VMware Virtualization Platform, on top of it Ansible Tower Platform and Automation of at least two OCP environments Deployment, each with its LB in Disconnected environment. In addition VMDK used as persistent storage backend for each OCP cluster and each cluster was with its own Metrics system (Prometheus/Grafana) and propagation of metrics to the Central Metrics environment.

Below the VMware vSphere details, Ansible Tower Details, Offline Registry and Repository Details

- VMware vSphere 6.7 environment
 - One Cluster
 - Andromeda01 & Andromeda02 physical Servers with ESXi 6.7 version
 - Andromeda02 - 4 Disks with 1.2 TB each one configured in RAID10 with total Storage 2.4TB defined as Datastore1 in vSphere
 - Andromeda01 - 2 Disks with 1.2 TB each one configured in RAID1 with total Storage 1.2TB defined as Datastore2 in vSphere, 2 Disks with 278 GB each one configured in RAID0 with total Storage 550 GB defined as Datastore3 in vSphere
 - HostedEngine VM appliance installed on top of ESXi 6.7 with DNS name of vSphere Manager: <https://vm-76-242.sales.lab.tlv.redhat.com>
 - 3 additional VMs:
 - 1 ocprepo (Offline Repository + Offline Container Registry + DNS)
 - 1 grafana (Central Grafana Server)
 - 1 anstower (Ansible Tower server)
- vSphere DataCenter details:
 - Datacenter name: Datasenter 1
 - Cluster name: Cluster 1
 - VM Folder names:
 - Common
 - OCP1
 - OCP2
- Red Hat Ansible Tower environment:
 - One node: VM name: anstower
 - Ansible Tower 3.4.3
 - Ansible 2.7.9
- Offline Repository + Offline Registry + DNS
 - Enabled Relevant Repos for OCP 3.11.98
 - Docker Distribution installed and all Images loaded, tagged and pushed for OCP 3.11.98
 - DNS installed and configured for ocp1.zayadtest.com and ocp2.zayadtest.com
- Central Grafana:
 - 1 VM - Grafana - with Grafana installed on it, version 6.1.4 and Prometheus as Data Source

3.1 Architecture

Here is a logical diagram of different parts of the implemented Ansible Tower Architecture.

vSphere & Ansible Tower Architecture

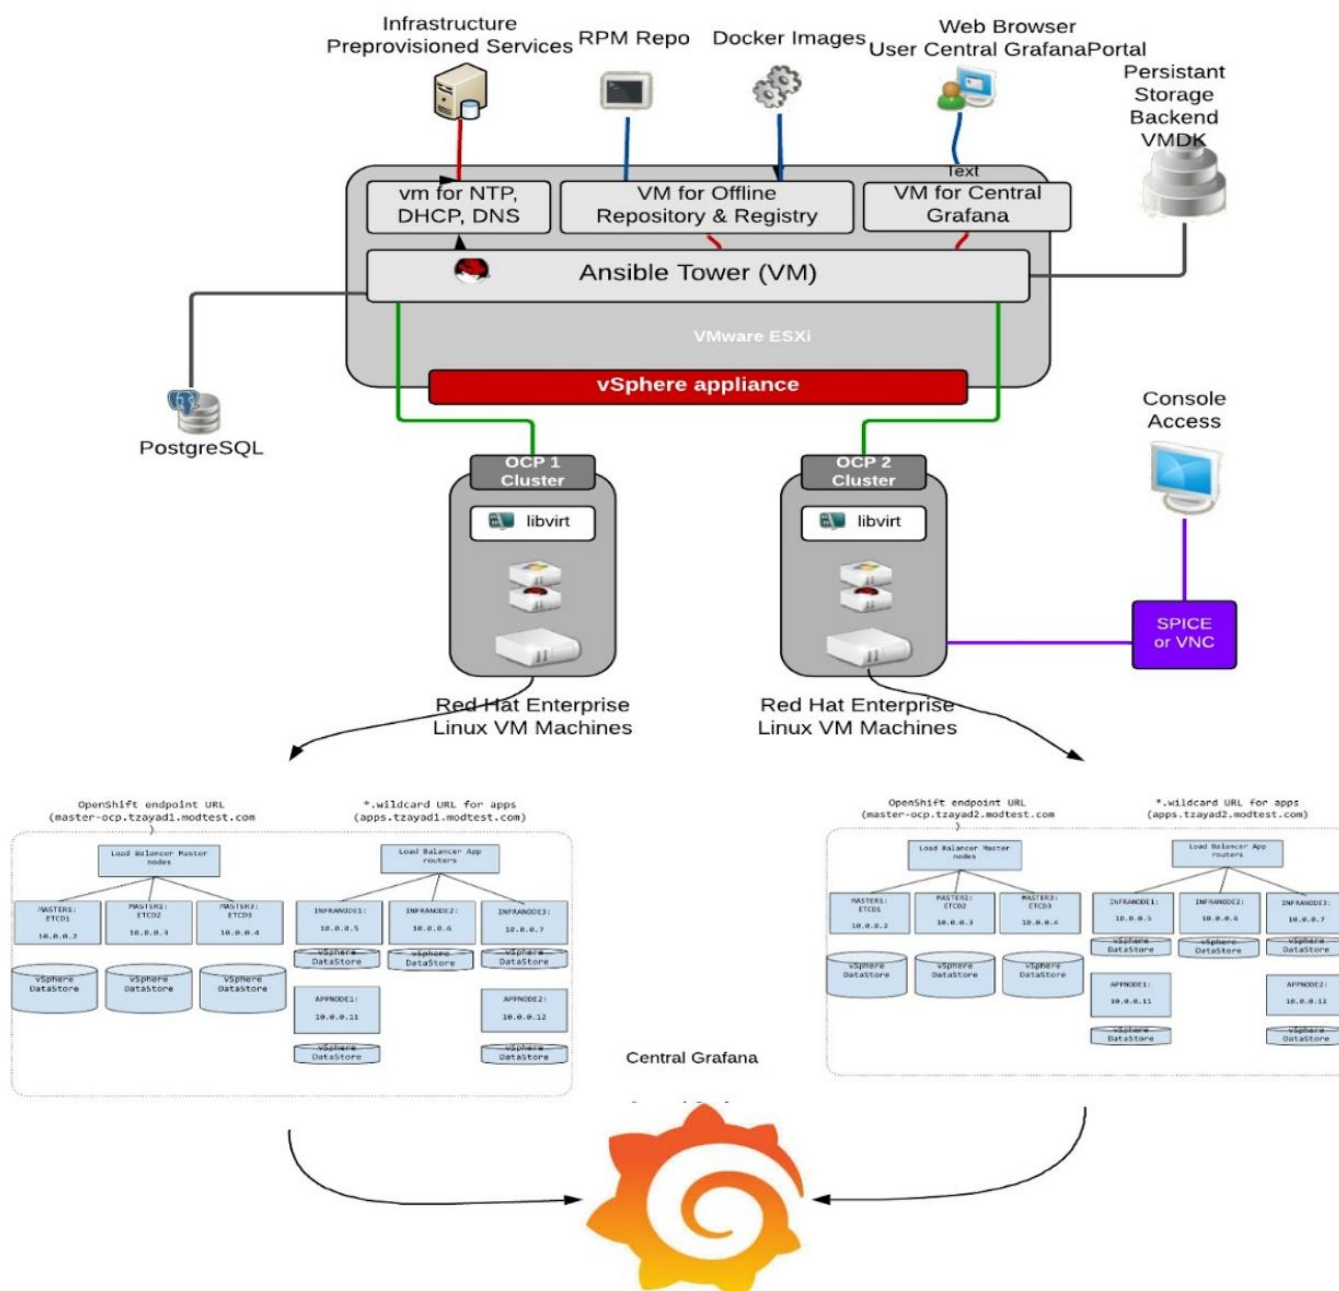


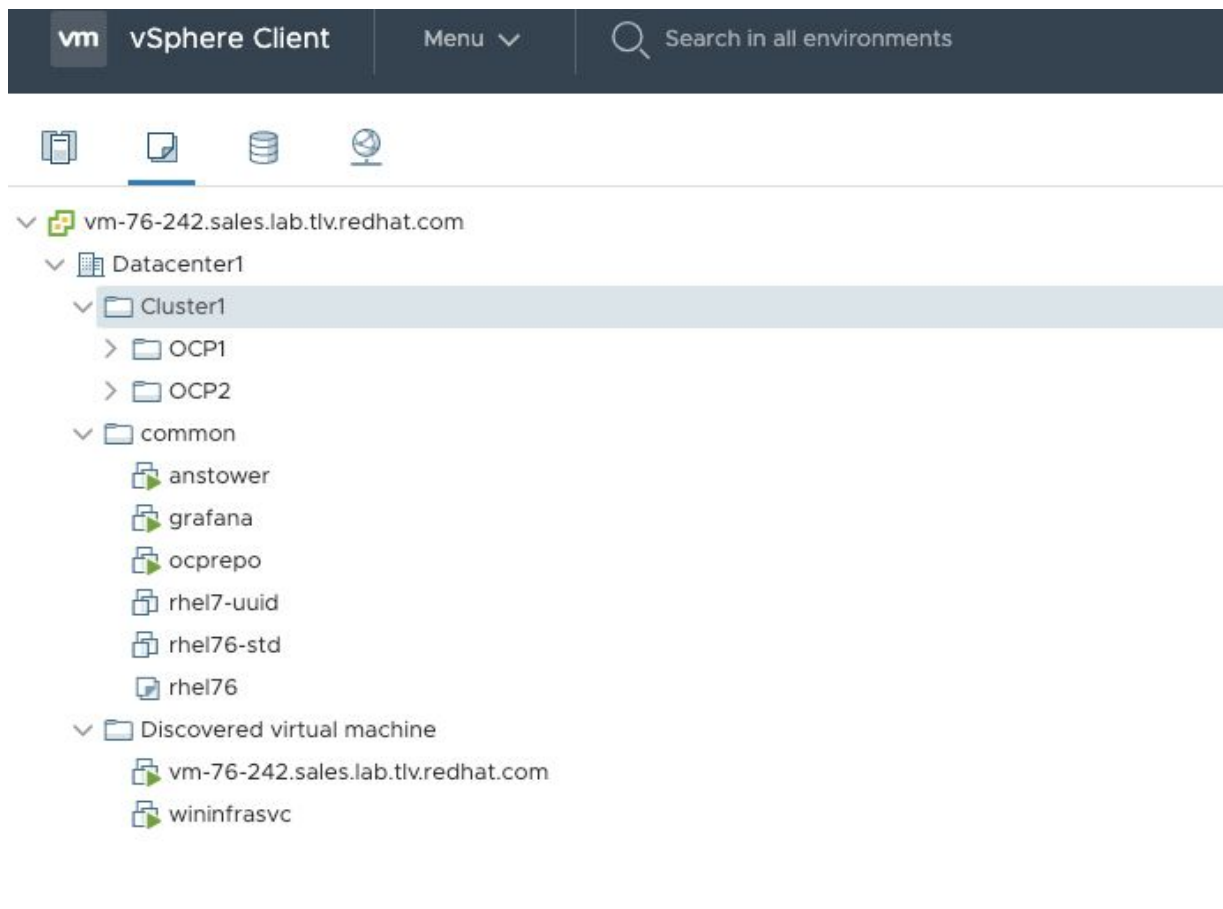
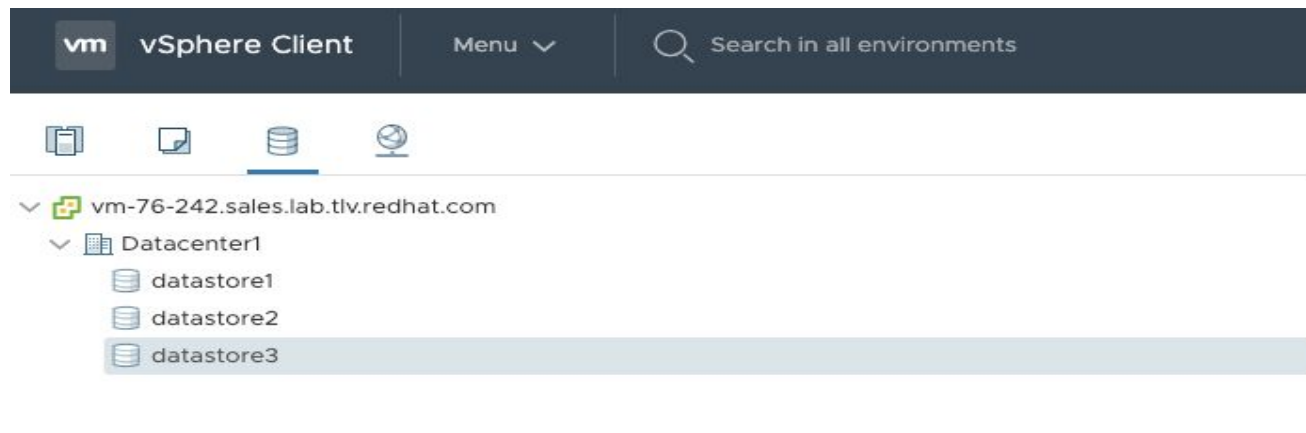
Image 3-2: <https://vm-76-242.sales.lab.tlv.redhat.com>- list of VMs.

Image 3-3: Logical Diagram of Datastores



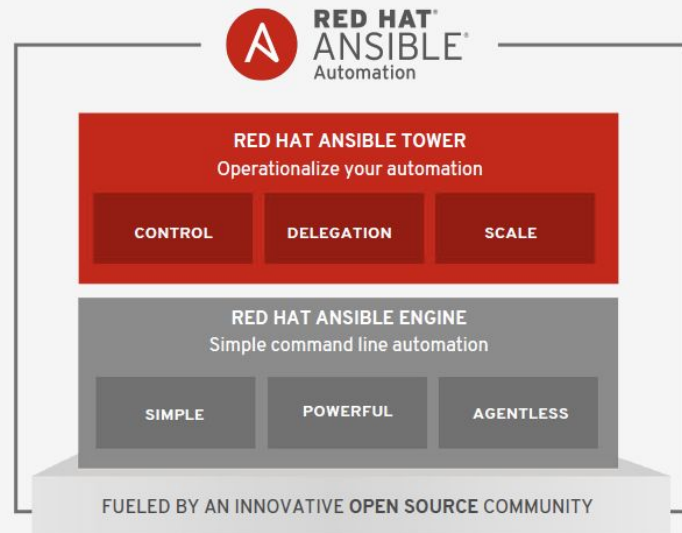
Ansible and Ansible Tower - Technical Overview

WHAT IS ANSIBLE AUTOMATION?

Ansible Automation is the enterprise **framework** for automating across IT operations.

Ansible Engine runs Ansible Playbooks, the automation **language** that can perfectly describes an IT application infrastructure.

Ansible Tower allows you **scale** IT automation, manage complex deployments and speed productivity.



WHY ANSIBLE?



SIMPLE

- Human readable automation
- No special coding skills needed
- Tasks executed in order
- Usable by every team
- Get productive quickly



POWERFUL

- App deployment
- Configuration management
- Workflow orchestration
- Network automation
- Orchestrate the app lifecycle



AGENTLESS

- Agentless architecture
- Uses OpenSSH & WinRM
- No agents to exploit or update
- Get started immediately
- More efficient & more secure

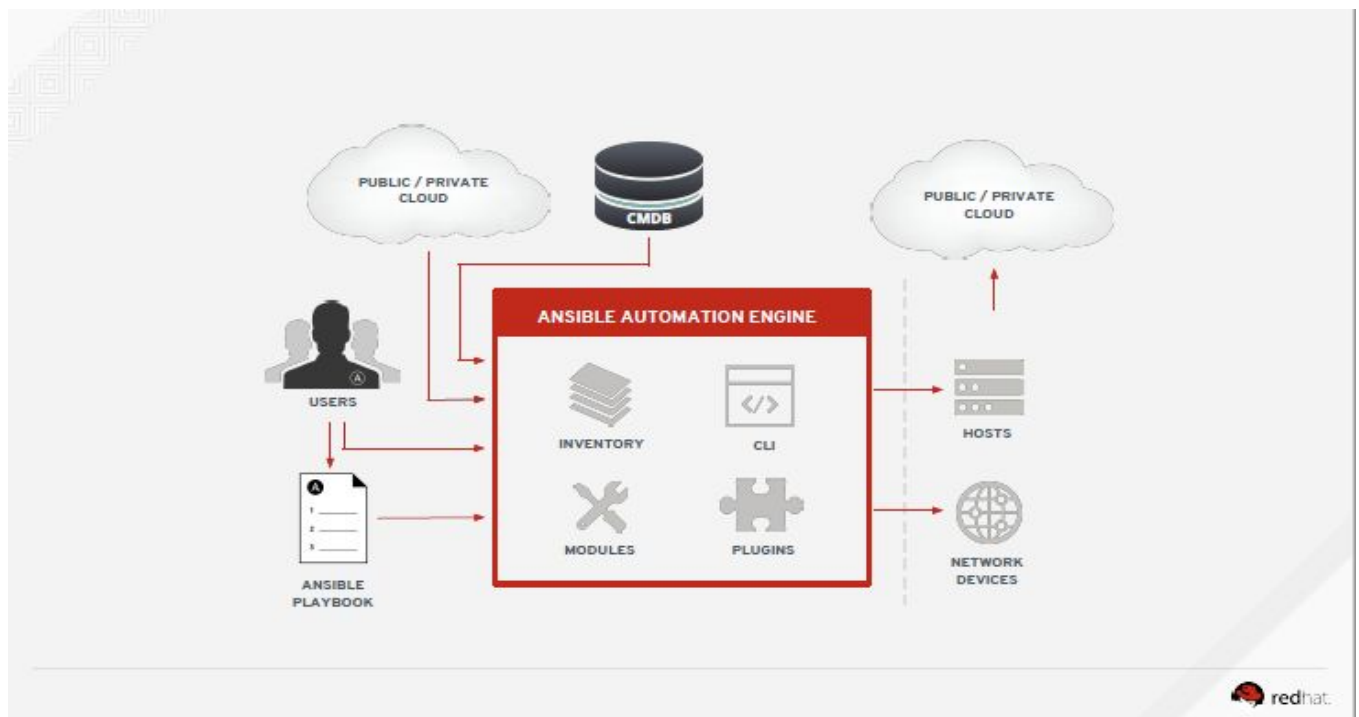
AUTOMATION FOR EVERYONE

```
---
- name: install and start apache
  hosts: web
  become: yes
  vars:
    http_port: 80

  tasks:
    - name: httpd package is present
      yum:
        name: httpd
        state: latest

    - name: latest index.html file is present
      copy:
        src: files/index.html
        dest: /var/www/html/

    - name: httpd is started
      service:
        name: httpd
        state: started
```



PLAYBOOK EXAMPLES:

GITHUB

github.com/ansible/ansible-examples

LAMP + HAProxy + NAGIOS

github.com/ansible/ansible-examples/tree/master/lamp_haproxy

WINDOWS

github.com/ansible/ansible-examples/tree/master/windows

SECURITY COMPLIANCE

github.com/ansible/ansible-lockdown

NETWORK AUTOMATION

ansible.com/linklight

github.com/network-automation



WHAT IS ANSIBLE TOWER?

Ansible Tower is a UI and RESTful API allowing you to scale IT automation, manage complex deployments and speed productivity.

- Role-based access control
- Deploy entire applications with push-button deployment access
- All automations are centrally logged
- Powerful workflows match your IT processes





RBAC

Allow restricting playbook access to authorized users. One team can use playbooks in check mode (read-only) while others have full administrative abilities.

PUSH BUTTON

An intuitive user interface experience makes it easy for novice users to execute playbooks you allow them access to.

RESTful API

With an API first mentality every feature and function of Tower can be API driven. Allow seamless integration with other tools like ServiceNow and Infoblox.

WORKFLOWS

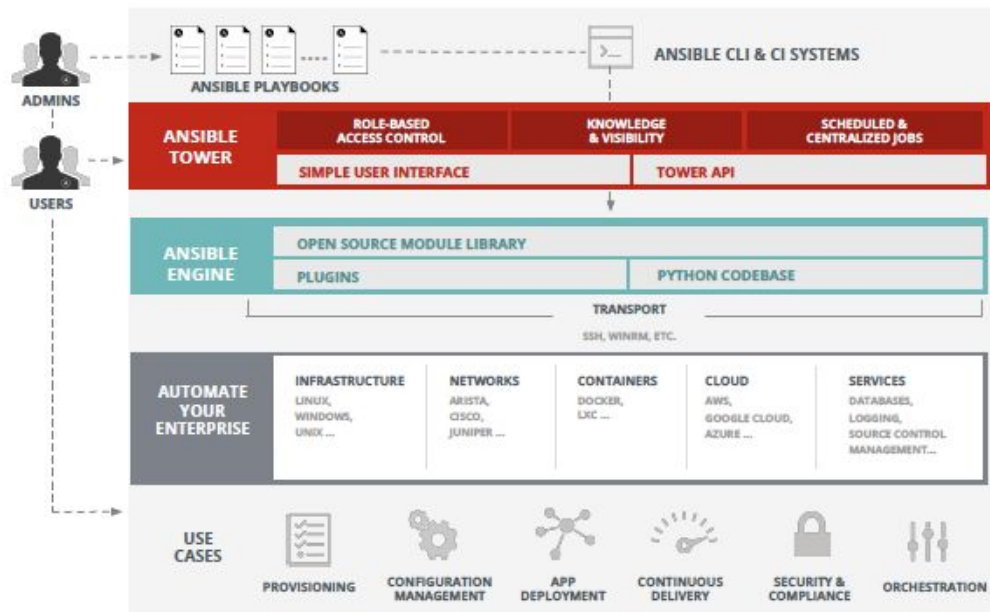
Ansible Tower's multi-playbook workflows chain any number of playbooks, regardless of whether they use different inventories, run as different users, run at once or utilize different credentials.

ENTERPRISE INTEGRATIONS

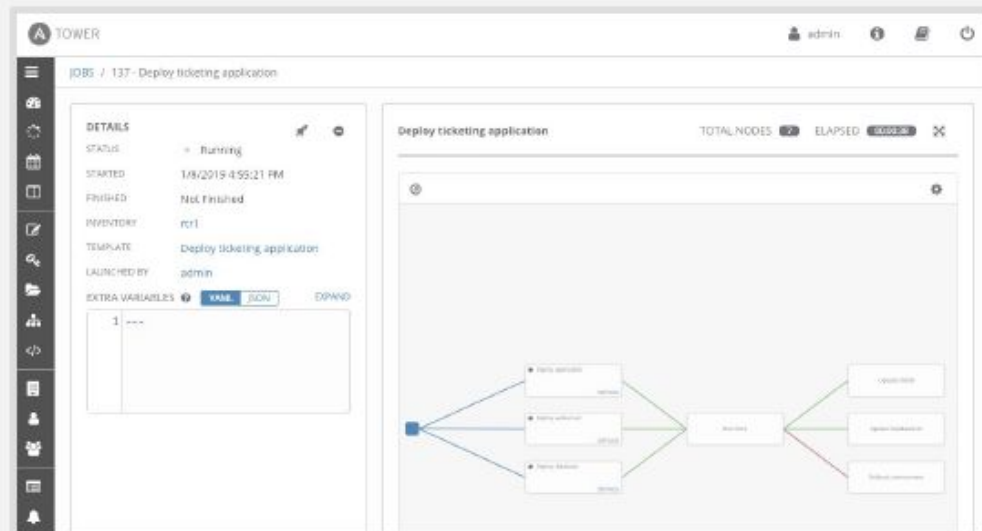
Integrate with enterprise authentication like TACACS+, RADIUS, Azure AD. Setup token authentication with OAuth 2. Setup notifications with PagerDuty, Slack and Twilio.

CENTRALIZED LOGGING

All automation activity is securely logged. Who ran it, how they customized it, what it did, where it happened - all securely stored and viewable later, or exported through Ansible Tower's API.



ANSIBLE TOWER FEATURES: CREATE AUTOMATION WORKFLOWS



NEXT STEPS

GET STARTED

ansible.com/get-started

ansible.com/tower-trial

JOIN THE COMMUNITY

ansible.com/community

WORKSHOPS & TRAINING

ansible.com/workshops

Red Hat Training

SHARE YOUR STORY

Follow us @Ansible

Friend us on Facebook



3.2 Hardware specifications

A. Hardware Detail

Can be found in the Ansible Tower Tzayad Raanana worksheet April 2019

[Ansible Tower Tzayad Raanana worksheet April 2019](#)

Table 3.2-1: Hardware details.

All the nodes in the cluster are Dell PowerEdge R430 with the next components:

Dell PowerEdge R430 * 1 (for vmware vsphere 6.7)

- 2x Intel(R) Xeon(R) CPU E5-2609 v4 @ 1.70GHz(Base) 8-core (16-thread) CPUs.
- 128GB RAM.
- 1x Integrated 1Gb quad port adapter.
- 1x Integrated RAID Controller
- RAID 1 1.2 TB *2 volume for esxi and the VMs that on top of him (Andromeda01)
- RAID 0 278 GB *2 volume for DataStore3 for pv (Andromeda01)

Dell PowerEdge R430 * 1 (for vmware vsphere 6.7)

- 2x Intel(R) Xeon(R) CPU E5-2609 v4 @ 1.70GHz(Base) 8-core (16-thread) CPUs.
- 128GB RAM.
- 1x Integrated 1Gb quad port adapter.
- 1x Integrated RAID Controller
- RAID 10 1.2 TB *4 volume for esxi and the VMs that on top of him(Andromeda2) and for pv

B. Network Details

| Network | Network | Netmask | Gateway | VLAN |
|------------|------------|-----------------|--------------|------|
| Vm network | 10.35.76.x | 255.255.255.128 | 10.35.76.254 | 461 |

Table 3-2: Network details.

C. Storage Details

Boot devices for physical hardware are provided by locally attached hardware RAID storage.

| Usage | Size | Type | Hostname | Additional |
|------------|--------|------|-------------------------|----------------|
| DataStore1 | 2.4 TB | scsi | Andromeda2 (not shared) | For VMs and pv |
| DataStore2 | 1.2 TB | scsi | Andromeda1 (not shared) | For VMs |
| DataStore3 | 577 GB | scsi | Andromeda1 (not shared) | For pv |

Table 3-3: Storage details.**D. Supporting Infrastructure**

- NTP Servers: 10.46.0.31
- DNS Servers:
 - 10.46.0.31
 - Search domain
 - sales.lab.tlv.redhat.com , zayadtest.com

E. Supporting Infrastructure

Here is the list of Software repositories that used on the install.

vmware vsphere 6.7 Infrastructure

Esxi 6.7 ISO, VMware-VCSA-all-6.7.0-11726888.iso, RHEL 7.6 ISO and Windows server 2016 ISO

Openshift 3.11 repositories

| Channel | Repository Name |
|--|--------------------------------|
| Red Hat Enterprise Linux 7 Server (RPMS) | rhel-7-server-rpms |
| Red Hat Enterprise Linux 7 Server - Extra (RPMS) | rhel-7-server-extras-rpms |
| Red Hat Enterprise ose 3.11 (RPMS) | rhel-7-server-ose-3.11-rpms |
| Red Hat Ansible Engine 2.6 RPMS | rhel-7-server-ansible-2.6-rpms |

Table 3-4: Required repositories.

4. ENGAGEMENT PHASES IN DETAILS

4.1 VMware vsphere 6.7 deployment

In this engagement the decision was to deploy VMware vsphere 6.7 as virtualization platform. This chapter outlines how to deploy VMware vsphere 6.7 across two servers. Deployment performed on servers Andromeda01, Andromeda02 - Management IP address 10.35.76.129, 10.35.76.130 according the relevant installation guide that can be found in the next link:

[Deploying VMware esxi 6.7](#)

Initially we deployed vcenter server appliance 6.7 on a windows server 2016 vm that we created on the esxi that we already installed. The relevant installation guide that can be found in the next link:

[Deploy vcenter server appliance 6.7](#)

The names of virtual machines and their roles can be found in the next link

[Ansible Tower Tzayad Raanana worksheet April 2019](#)

4.2 Offline Repository deployment

The following procedure is for preparing your own offline deployer. The procedure is based on the following document:

[Enabling the ose 3.11 Repositories standalone install](#)

Setup when using online server (We used GCP instance already provisioned for this purpose)

Goal:

Obtaining required software packages and images

Before you install OpenShift Container Platform in your disconnected environment, obtain the required images and components and store them in your repository.

You must obtain the required images and software components on a system with the same architecture as the cluster that is in your disconnected environment

For this purpose Virtual Instance was created on Red Hat PS GCP environment , (instance name *ocp-offline-repo-registry*), you can ask permission to use it if you need (send request to Anael with your id_rsa.public key)

Obtaining OpenShift Container Platform packages

On the RHEL 7 server with an internet connection, sync the repositories:

1. To ensure that the packages are not deleted after you sync the repository, import the GPG key:

```
$ rpm --import /etc/pki/rpm-gpg/RPM-GPG-KEY-redhat-release
```

2. If **subscription-manager** component not installed please install it:

```
$ sudo yum install subscription-manager
```

3. Register the server with the Red Hat Customer Portal. You must use the credentials that are associated with the account that has access to the OpenShift Container Platform subscriptions:

```
$ sudo subscription-manager register
```

4. Pull the latest subscription data from RHSM:

```
$ sudo subscription-manager refresh
```

5. Attach a pool ID for a subscription that provides OpenShift Container Platform:

```
$ sudo subscription-manager attach --pool=<pool_id>
```

```
$ sudo subscription-manager repos --disable="*"
```

6. Enable only the repositories required by OpenShift Container Platform 3.11:

```
$ sudo subscription-manager repos \
    --enable="rhel-7-server-rpms" \
    --enable="rhel-7-server-extras-rpms" \
    --enable="rhel-7-server-ose-3.11-rpms" \
    --enable="rhel-7-server-ansible-2.6-rpms"
```

7. Install required packages:

```
$ sudo yum -y install yum-utils createrepo docker mlocate screen git
$ sudo updatedb
```

The **yum-utils** package provides the **reposync** utility, which lets you mirror yum repositories, and you can use the **createrepo** package to create a usable **yum** repository from a directory.

8. Make a directory to store the software in the server's storage or to a USB drive or other external device:

```
$ sudo mkdir -p </path/to/repos>
```

If you can re-connect this server to the disconnected LAN and use it as the repository server, store the files locally. If you cannot, use USB-connected storage so you can transport the software to a repository server in your disconnected LAN

9. Sync the packages and create the repository for each of them
 - o For on-premise installations on x86_64 servers, run the following command

```
$ sudo for repo in \
    rhel-7-server-rpms \
    rhel-7-server-extras-rpms \
    rhel-7-server-ansible-2.6-rpms \
    rhel-7-server-ose-3.11-rpms
```

do

```
reposync --gpgcheck -lm --repoid=${repo}
```

```
--download_path=</path/to/repos>
```

```
createrepo -v </path/to/repos/>${repo} -o </path/to/repos/>${repo}
```

Done

Obtaining images.

Pull the required container images:

1. Start the Docker daemon:

```
$ sudo systemctl start docker
```

2. Pull all of the required OpenShift Container Platform infrastructure component images. Replace <tag> with the version to install. For example, specify v3.11.98 for the latest version. You can specify a different minor version.

Pull all of the required OpenShift Container Platform component images for the optional components. Replace <tag> with the version to install. For example, specify v3.11.98 for the latest version. You can specify a different minor version.

Pull the Red Hat-certified [Source-to-Image \(S2I\)](#) builder images that you intend to use in your OpenShift Container Platform environment.

- We created script that performs pull of all relevant docker images -

pull_infrastructure_component_images.sh

```
#!/bin/bash
for image in \
apb-base:v3.11.98 \
apb-tools:v3.11.98 \
automation-broker-apb:v3.11.98 \
csi-attacher:v3.11.98 \
csi-driver-registrar:v3.11.98 \
csi-livenessprobe:v3.11.98 \
csi-provisioner:v3.11.98 \
grafana:v3.11.98 \
image-inspector:v3.11.98 \
local-storage-provisioner:v3.11.98 \
manila-provisioner:v3.11.98 \
mariadb-apb:v3.11.98 \
mediawiki:v3.11.98 \
mediawiki-apb:v3.11.98 \
```

```
mysql-apb:v3.11.98 \  
ose-ansible:v3.11.98 \  
ose-ansible-service-broker:v3.11.98 \  
ose-cli:v3.11.98 \  
ose-cluster-autoscaler:v3.11.98 \  
ose-cluster-capacity:v3.11.98 \  
ose-cluster-monitoring-operator:v3.11.98 \  
ose-console:v3.11.98 \  
ose-configmap-reloader:v3.11.98 \  
ose-control-plane:v3.11.98 \  
ose-deployer:v3.11.98 \  
ose-descheduler:v3.11.98 \  
ose-docker-builder:v3.11.98 \  
ose-docker-registry:v3.11.98 \  
ose-efs-provisioner:v3.11.98 \  
ose-egress-dns-proxy:v3.11.98 \  
ose-egress-http-proxy:v3.11.98 \  
ose-egress-router:v3.11.98 \  
ose-haproxy-router:v3.11.98 \  
ose-hyperkube:v3.11.98 \  
ose-hypershift:v3.11.98 \  
ose-keepalived-ipfailover:v3.11.98 \  
ose-kube-rbac-proxy:v3.11.98 \  
ose-kube-state-metrics:v3.11.98 \  
ose-metrics-server:v3.11.98 \  
ose-node:v3.11.98 \  
ose-node-problem-detector:v3.11.98 \  
ose-operator-lifecycle-manager:v3.11.98 \  
ose-ovn-kubernetes:v3.11.98 \  
ose-pod:v3.11.98 \  
ose-prometheus-config-reloader:v3.11.98 \  
ose-prometheus-operator:v3.11.98 \  
ose-recycler:v3.11.98 \  
ose-service-catalog:v3.11.98 \  
ose-template-service-broker:v3.11.98 \  
ose-tests:v3.11.98 \  
ose-web-console:v3.11.98 \  
postgresql-apb:v3.11.98 \  
registry-console:v3.11.98 \  
snapshot-controller:v3.11.98 \  
snapshot-provisioner:v3.11.98; \  
do \  
    sudo docker pull registry.access.redhat.com/openshift3/$image  
done
```

```
sudo docker pull registry.access.redhat.com/rhel7/etcd:3.2.22
```

```
for image in \  
metrics-cassandra:v3.11.98 \  
metrics-hawkular-metrics:v3.11.98 \  
metrics-hawkular-openshift-agent:v3.11.98 \  
metrics-heapster:v3.11.98 \  
metrics-schema-installer:v3.11.98 \  
oauth-proxy:v3.11.98 \  
ose-logging-curator5:v3.11.98 \  
ose-logging-elasticsearch5:v3.11.98 \  
ose-logging-eventrouter:v3.11.98 \  
ose-logging-fluentd:v3.11.98 \  
ose-logging-kibana5:v3.11.98 \  
prometheus:v3.11.98 \  
prometheus-alert-buffer:v3.11.98 \  
prometheus-alertmanager:v3.11.98 \  
prometheus-node-exporter:v3.11.98; \  
do \  
    sudo docker pull registry.access.redhat.com/openshift3/$image  
done
```

```
for image in \  
cfme-openshift-postgresql \  
cfme-openshift-memcached \  
cfme-openshift-app-ui \  
cfme-openshift-app \  
cfme-openshift-embedded-ansible \  
cfme-openshift-httpd \  
cfme-httpd-configmap-generator; \  
do \  
    sudo docker pull registry.access.redhat.com/cloudforms46/$image  
done
```

```
for image in \  
rhgs-server-rhel7 \  
rhgs-volmanager-rhel7 \  
rhgs-gluster-block-prov-rhel7 \  
rhgs-s3-server-rhel7; \  
do \  
    sudo docker pull registry.access.redhat.com/rhgs3/$image  
done
```

```
for image in \  
jboss-amq-6/amq63-openshift \  
done
```

```
jboss-datagrid-7/datagrid71-openshift \  
jboss-datagrid-7/datagrid71-client-openshift \  
jboss-datavirt-6/datavirt63-openshift \  
jboss-datavirt-6/datavirt63-driver-openshift \  
jboss-decisionserver-6/decisionserver64-openshift \  
jboss-processserver-6/processserver64-openshift \  
jboss-eap-6/eap64-openshift \  
jboss-eap-7/eap70-openshift \  
jboss-webserver-3/webserver31-tomcat7-openshift \  
jboss-webserver-3/webserver31-tomcat8-openshift \  
openshift3/jenkins-2-rhel7 \  
openshift3/jenkins-agent-maven-35-rhel7 \  
openshift3/jenkins-agent-nodejs-8-rhel7 \  
openshift3/jenkins-slave-base-rhel7 \  
openshift3/jenkins-slave-maven-rhel7 \  
openshift3/jenkins-slave-nodejs-rhel7 \  
rhsc1/mongodb-32-rhel7 \  
rhsc1/mysql-57-rhel7 \  
rhsc1/perl-524-rhel7 \  
rhsc1/php-56-rhel7 \  
rhsc1/postgresql-95-rhel7 \  
rhsc1/python-35-rhel7 \  
redhat-sso-7/sso70-openshift \  
rhsc1/ruby-24-rhel7 \  
redhat-openjdk-18/openjdk18-openshift \  
redhat-sso-7/sso71-openshift \  
rhsc1/nodejs-6-rhel7 \  
rhsc1/mariadb-101-rhel7; \  
do \  
    sudo docker pull registry.access.redhat.com/$image  
done
```

Exporting images

Since your environment does not have access to your internal network and requires physical media to transfer content, export the images to compressed files.

1. Create a directory to store your compressed images in and change to it:

```
$ mkdir </path/to/images>  
$ cd </path/to/images>
```

2. Export the OpenShift Container Platform infrastructure component images:
 - a. For on-premise installations on x86_64 servers, run the following command:

```
$ sudo docker save -o ose3-images.tar \  

```

```
registry.redhat.io/openshift3/apb-base \  
registry.redhat.io/openshift3/apb-tools \  
registry.redhat.io/openshift3/automation-broker-apb \  
registry.redhat.io/openshift3/csi-attacher \  
registry.redhat.io/openshift3/csi-driver-registrar \  
registry.redhat.io/openshift3/csi-livenessprobe \  
registry.redhat.io/openshift3/csi-provisioner \  
registry.redhat.io/openshift3/grafana \  
registry.redhat.io/openshift3/image-inspector \  
registry.redhat.io/openshift3/local-storage-provisioner \  
registry.redhat.io/openshift3/manila-provisioner \  
registry.redhat.io/openshift3/mariadb-apb \  
registry.redhat.io/openshift3/mediawiki \  
registry.redhat.io/openshift3/mediawiki-apb \  
registry.redhat.io/openshift3/mysql-apb \  
registry.redhat.io/openshift3/ose-ansible \  
registry.redhat.io/openshift3/ose-ansible-service-broker \  
registry.redhat.io/openshift3/ose-cli \  
registry.redhat.io/openshift3/ose-cluster-autoscaler \  
registry.redhat.io/openshift3/ose-cluster-capacity \  
registry.redhat.io/openshift3/ose-cluster-monitoring-operator \  
registry.redhat.io/openshift3/ose-console \  
registry.redhat.io/openshift3/ose-configmap-reloader \  
registry.redhat.io/openshift3/ose-control-plane \  
registry.redhat.io/openshift3/ose-deployer \  
registry.redhat.io/openshift3/ose-descheduler \  
registry.redhat.io/openshift3/ose-docker-builder \  
registry.redhat.io/openshift3/ose-docker-registry \  
registry.redhat.io/openshift3/ose-efs-provisioner \  
registry.redhat.io/openshift3/ose-egress-dns-proxy \  
registry.redhat.io/openshift3/ose-egress-http-proxy \  
registry.redhat.io/openshift3/ose-egress-router \  
registry.redhat.io/openshift3/ose-haproxy-router \  
registry.redhat.io/openshift3/ose-hyperkube \  
registry.redhat.io/openshift3/ose-hypershift \  
registry.redhat.io/openshift3/ose-keepalived-ipfailover \  
registry.redhat.io/openshift3/ose-kube-rbac-proxy \  
registry.redhat.io/openshift3/ose-kube-state-metrics \  
registry.redhat.io/openshift3/ose-metrics-server \  
registry.redhat.io/openshift3/ose-node \  
registry.redhat.io/openshift3/ose-node-problem-detector \  
registry.redhat.io/openshift3/ose-operator-lifecycle-manager \  
registry.redhat.io/openshift3/ose-ovn-kubernetes \  
registry.redhat.io/openshift3/ose-pod \  
registry.redhat.io/openshift3/ose-prometheus-config-reloader \  

```



```
registry.redhat.io/openshift3/ose-prometheus-operator \  
registry.redhat.io/openshift3/ose-recycler \  
registry.redhat.io/openshift3/ose-service-catalog \  
registry.redhat.io/openshift3/ose-template-service-broker \  
registry.redhat.io/openshift3/ose-tests \  
registry.redhat.io/openshift3/ose-web-console \  
registry.redhat.io/openshift3/postgresql-apb \  
registry.redhat.io/openshift3/registry-console \  
registry.redhat.io/openshift3/snapshot-controller \  
registry.redhat.io/openshift3/snapshot-provisioner \  
registry.redhat.io/rhel7/etcd
```

- b. For optional components, to export them, run the following command:

```
$ sudo docker save -o ose3-optional-imgs.tar \  
registry.redhat.io/openshift3/metrics-cassandra \  
registry.redhat.io/openshift3/metrics-hawkular-metrics \  
registry.redhat.io/openshift3/metrics-hawkular-openshift-agent \  
registry.redhat.io/openshift3/metrics-heapster \  
registry.redhat.io/openshift3/metrics-schema-installer \  
registry.redhat.io/openshift3/oauth-proxy \  
registry.redhat.io/openshift3/ose-logging-curator5 \  
registry.redhat.io/openshift3/ose-logging-elasticsearch5 \  
registry.redhat.io/openshift3/ose-logging-eventrouter \  
registry.redhat.io/openshift3/ose-logging-fluentd \  
registry.redhat.io/openshift3/ose-logging-kibana5 \  
registry.redhat.io/openshift3/prometheus \  
registry.redhat.io/openshift3/prometheus-alert-buffer \  
registry.redhat.io/openshift3/prometheus-alertmanager \  
registry.redhat.io/openshift3/prometheus-node-exporter \  
registry.redhat.io/cloudforms46/cfme-openshift-postgresql \  
registry.redhat.io/cloudforms46/cfme-openshift-memcached \  
registry.redhat.io/cloudforms46/cfme-openshift-app-ui \  
registry.redhat.io/cloudforms46/cfme-openshift-app \  
registry.redhat.io/cloudforms46/cfme-openshift-embedded-ansible \  
\  
registry.redhat.io/cloudforms46/cfme-openshift-httpd \  
registry.redhat.io/cloudforms46/cfme-httpd-configmap-generator \  
registry.redhat.io/rhgs3/rhgs-server-rhel7 \  
registry.redhat.io/rhgs3/rhgs-volmanager-rhel7 \  
registry.redhat.io/rhgs3/rhgs-gluster-block-prov-rhel7 \  
registry.redhat.io/rhgs3/rhgs-s3-server-rhel7
```

- c. Export the S2I builder images that you pulled.

```
sudo docker save -o ose3-builder-images.tar \
```

```
registry.access.redhat.com/jboss-webserver-3/webserver31-tomcat7-openshift \
registry.access.redhat.com/jboss-amq-6/amq63-openshift \
registry.access.redhat.com/jboss-datagrid-7/datagrid71-openshift \
registry.access.redhat.com/jboss-datagrid-7/datagrid71-client-openshift \
registry.access.redhat.com/jboss-datavirt-6/datavirt63-openshift \
registry.access.redhat.com/jboss-datavirt-6/datavirt63-driver-openshift \
```

```
registry.access.redhat.com/jboss-decisionserver-6/decisionserver64-openshift \
```

```
registry.access.redhat.com/jboss-processserver-6/processserver64-openshift \
registry.access.redhat.com/jboss-eap-6/eap64-openshift \
registry.access.redhat.com/jboss-eap-7/eap70-openshift \
```

```
registry.access.redhat.com/jboss-webserver-3/webserver31-tomcat7-openshift \
```

```
registry.access.redhat.com/jboss-webserver-3/webserver31-tomcat8-openshift \
registry.access.redhat.com/openshift3/jenkins-2-rhel7 \
registry.access.redhat.com/openshift3/jenkins-agent-maven-35-rhel7 \
registry.access.redhat.com/openshift3/jenkins-agent-nodejs-8-rhel7 \
registry.access.redhat.com/openshift3/jenkins-slave-base-rhel7 \
registry.access.redhat.com/openshift3/jenkins-slave-maven-rhel7 \
registry.access.redhat.com/openshift3/jenkins-slave-nodejs-rhel7 \
registry.access.redhat.com/rhsc1/mongodb-32-rhel7 \
registry.access.redhat.com/rhsc1/mysql-57-rhel7 \
registry.access.redhat.com/rhsc1/perl-524-rhel7 \
registry.access.redhat.com/rhsc1/php-56-rhel7 \
registry.access.redhat.com/rhsc1/postgresql-95-rhel7 \
registry.access.redhat.com/rhsc1/python-35-rhel7 \
registry.access.redhat.com/redhat-sso-7/sso70-openshift \
registry.access.redhat.com/rhsc1/ruby-24-rhel7 \
registry.access.redhat.com/redhat-openjdk-18/openjdk18-openshift \
registry.access.redhat.com/redhat-sso-7/sso71-openshift \
registry.access.redhat.com/rhsc1/nodejs-6-rhel7 \
registry.access.redhat.com/rhsc1/mariadb-101-rhel7
```

d. Copy the compressed files from your Internet-connected host to your internal host.

```
-rw-----, 1 root root 12G Dec 20 13:00 ose3-images.tar
-rw-----, 1 root root 7.3G Dec 20 14:29 ose3-optional-images.tar
-rw-----, 1 root root 11G Dec 20 14:53 ose3-builder-images.tar
```

Exporting repos

Since your environment does not have access to your internal network and requires physical media to transfer content, export the repos to compressed files.

1. Create a directory to store your compressed images in and change to it:

```
$ mkdir </path/to/repos>  
$ cd </path/to/repos>
```

2. Export the OpenShift Container Platform infrastructure component repos:

```
$ sudo tar -czvf ocp3.11.98.repo.tgz ocprepo/
```

3. Copy the compressed files from your Internet-connected host to your internal host.

```
-rw-r--r--. 1 root root 36G Dec 20 13:57 ocp3.11.98.repo.tgz
```

Setup on internal (not connected to the internet) VM: ocprepo

1. Deploy instance using RHEL 7.6 image
Note: httpd should already be installed on the image
2. Copy prepackaged tar files to /var/www/html/
3. Unarchive the repos in the /var/www/html/repos
4. Check mode on repos directory
chmod -R 7777 repos
5. Open firewall port 80/tcp
systemctl status firewalld.service
systemctl start firewalld.service
systemctl enable firewalld.service
firewall-cmd --permanent --add-port=80/tcp
firewall-cmd --reload
6. Enable and Start httpd service

4.3 Offline Container registry deployment

This section covers the setup of an offline container registry. The procedure is based on the installation guide [Creating an Offline Container Registry for openshift 3.11](#)

Setup when using offline server

1. Deploy instance using RHEL 7.6 image (alternatively use the same instance as in section 4.2)
2. Install local registry
yum docker-distribution -y
systemctl enable docker-distribution
systemctl start docker-distribution
3. Edit /etc/docker/daemon.json so it has the local registry as insecure:
{
 "insecure-registries" : ["<registry/p>:5000"]
}

- ```
}
4. Restart docker process
5. Re-Tag images
 Run the following command for each image:
 note: <tag> value must match the value defined in local_registry_images.yaml on directory. In this
 deployment tags are 'v3.11.98', 'v3.11', 'latest':
 docker tag <registryIp>:5000/<image>:<tag>
6. Push to local registry
 docker push <image>:<tag>
9. Open firewall port 5000/tcp
 systemctl status firewalld.service
 systemctl start firewalld.service
 systemctl enable firewalld.service
 firewall-cmd --permanent --add-port=5000/tcp
 firewall-cmd --reload
```

## 4.4 Ansible Tower

At the start we downloaded the tar file from this website: [ansible tower tar file](#).

20 GB of dedicated hard disk space for Tower service nodes and for nodes containing a database 150 GB+ recommended. Mount this disks to /var in the node.

After this we edited the inventory file to our needs.

Then we run the ./setup.sh script and it installed ansible tower.

Edit the /etc/ansible/ansible.cfg forks = <number that you want> (default is 5, we increased it to 100)

Initially we need to install govc on the ansible tower server: [guide to install GOVC](#).

For backup run the ./setup.sh -b script.

After installation License required to be able to complete initial configuration. We used

[https://store.ansible.com/redhat/tower\\_license/](https://store.ansible.com/redhat/tower_license/) - Ansible Enterprise license generator FOR RED HATTERS (need valid Red Hat email address - for internal use only)

## 4.5 central grafana

Follow these steps to install central grafana: [central grafana](#)

## 4.6 Template creation in vSphere for OCP VMs

First we created in vSphere rhel 7.6 template that contains a minimal installation with some addons to it:


- We enabled the disk uuid option.
- Disabled the swap.
- Enabled Network Manager service.
- Enabled SELinux in enforced mode.
- Enabled the IP forwarding.

- We entered the public key of the ansible tower server in the authorized\_keys file.
- Created a private key of the root user of the ansible tower, because this is necessary for the ansible tower passwordless ssh access to managed hosts. When we start a job we need to choose which private key we want to use.

## 4.7 Ansible Tower configurations and definitions


We created three projects:

- The first project was ocpprep which contains the content of a git repository [ocpprep repository](#):




The screenshot shows the Ansible Tower web interface in a Google Chrome browser window. The browser's address bar shows the URL `https://10.35.76.250/#/projects/10`. The interface has a dark sidebar on the left with a menu containing 'Dashboard', 'Jobs', 'Schedules', 'My View', 'Templates', 'Credentials', 'Projects', 'Inventories', 'Inventory Scripts', 'Organizations', 'Users', 'Teams', and 'Administration'. The main content area is titled 'PROJECTS / ocpprep' and contains a form for configuring the 'ocpprep' project. The form has tabs for 'DETAILS', 'PERMISSIONS', 'NOTIFICATIONS', 'JOB TEMPLATES', and 'SCHEDULES'. The 'DETAILS' tab is active, showing fields for 'NAME' (ocpprep), 'DESCRIPTION' (prep and post playbooks), 'ORGANIZATION' (Default), 'SCM TYPE' (Git), 'SCM URL' (https://github.com/rockocoop/ocpprep.git), 'SCM BRANCH/TAG/COMMIT', 'SCM CREDENTIAL', 'SCM UPDATE OPTIONS' (Clean, Delete on Update, Update Revision on Launch), and 'CACHE TIMEOUT (SECONDS)' (0). At the bottom right of the form are 'CANCEL' and 'SAVE' buttons. The bottom of the browser window shows a terminal prompt `event@event:~` and a tab for 'Ansible Tower | ocpprep - Google ...'.

- The second project was openshift-ansible which contains all the installation playbooks and it's located locally on the tower server under `/var/lib/awx/projects` folder:



The screenshot shows the Ansible Tower web interface in Google Chrome. The browser tabs include 'Ansible Tower | openshift-ansible', 'Inbox (2,426) - matan.carr...', 'Ansible Tower MOD Tzayad', and 'GitHub - rockocoop/ocpp...'. The address bar shows 'https://10.35.76.250/#/projects/12'. The left sidebar has a menu with 'TOWER' at the top, followed by 'VIEWS' (Dashboard, Jobs, Schedules, My View), 'RESOURCES' (Templates, Credentials, Projects, Inventories, Inventory Scripts), 'ACCESS' (Organizations, Users, Teams), and 'ADMINISTRATION' (Credential Types). The main content area is titled 'PROJECTS / openshift-ansible'. It has tabs for 'DETAILS', 'PERMISSIONS', 'NOTIFICATIONS', 'JOB TEMPLATES', and 'SCHEDULES'. The 'DETAILS' tab is active, showing fields for NAME (openshift-ansible), DESCRIPTION, ORGANIZATION (Default), SCM TYPE (Manual), PROJECT BASE PATH (/var/lib/awx/projects), and PLAYBOOK DIRECTORY (openshift-ansible). Below these fields is a 'PROJECTS' section with a search bar and a table listing projects. The table has columns for project name, revision, organization, last modified, and last used. The first project listed is 'ocpprep' with revision '6b166a1', organization 'Default', last modified '4/29/2019 12:29:01 PM', and last used '4/29/2019 12:29:01 PM'. The bottom of the browser shows a terminal window with 'event@event:~' and a taskbar with 'Ansible Tower | openshift-ansible'.

- The third project was vmware-ansible which contained the playbooks that created the VMs of the OCP cluster or destroy them [vmware-ansible](#):



The screenshot shows the Ansible Tower web interface in Google Chrome. The browser tabs include 'Ansible Tower | vmware-ansible', 'Inbox (2,426) - matan.carr...', 'Ansible Tower MOD Tzayad', and 'GitHub - rockocoop/ocpp...'. The address bar shows 'https://10.35.76.250/#/projects/6'. The left sidebar is identical to the previous screenshot. The main content area is titled 'PROJECTS / vmware-ansible'. It has tabs for 'DETAILS', 'PERMISSIONS', 'NOTIFICATIONS', 'JOB TEMPLATES', and 'SCHEDULES'. The 'DETAILS' tab is active, showing fields for NAME (vmware-ansible), DESCRIPTION, ORGANIZATION (Default), SCM TYPE (Git), and SOURCE DETAILS. The SOURCE DETAILS section includes SCM URL (https://github.com/rockocoop/vmware-ansible.git), SCM BRANCH/TAG/COMMIT, SCM CREDENTIAL, and SCM UPDATE OPTIONS (Clean, Delete on Update, Update Revision on Launch). Below these fields is a 'PROJECTS' section with a search bar and a table listing projects. The table has columns for project name, revision, organization, last modified, and last used. The first project listed is 'ocpprep' with revision '6b166a1', organization 'Default', last modified '4/29/2019 12:29:01 PM', and last used '4/29/2019 12:29:01 PM'. The bottom of the browser shows a terminal window with 'event@event:~' and a taskbar with 'Ansible Tower | vmware-ansible'.

We also created two more project that looked like the first and the third projects but the difference was that the playbooks were located locally on the tower server instead of being on git and we tested it and it worked.

After that we created inventory for the jobs. The first section of the inventory is the details, it contains variables on all of the hosts:

---

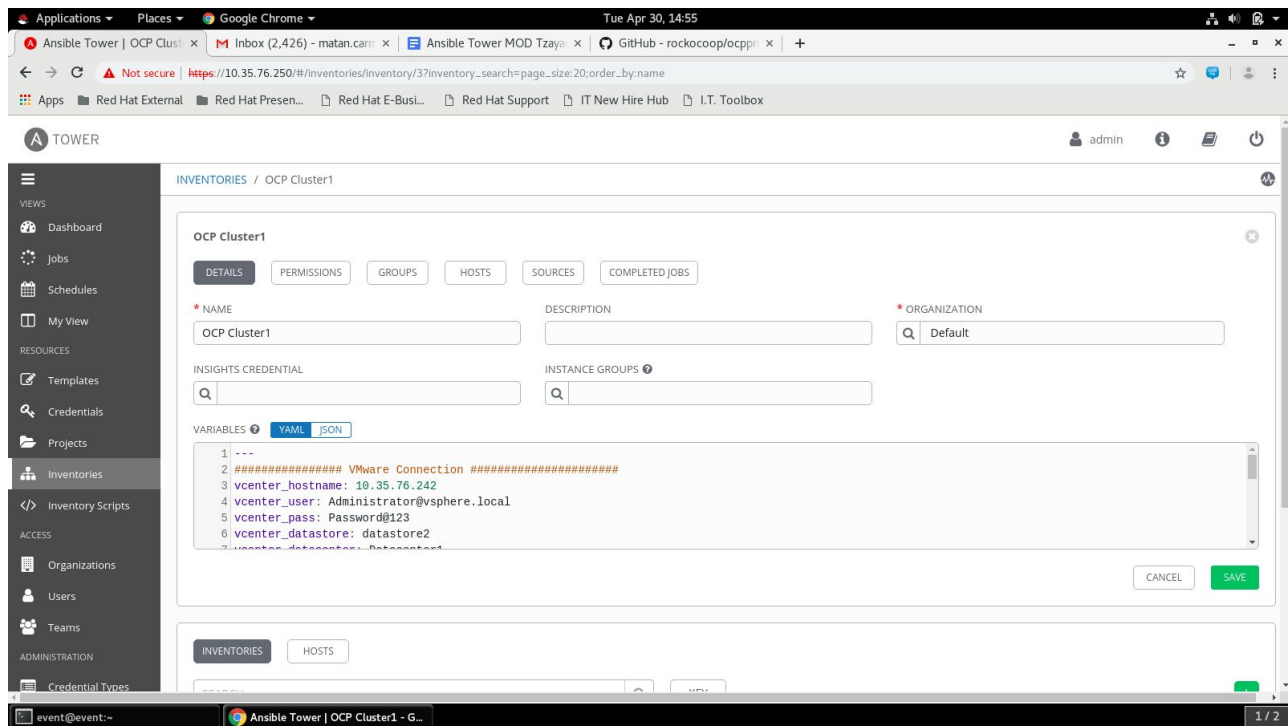
##### VMware Connection #####

vcenter\_hostname: 10.35.76.242  
vcenter\_user: Administrator@vsphere.local  
vcenter\_pass: Password@123  
vcenter\_datastore: datastore2  
vcenter\_datacenter: Datacenter1  
vcenter\_cluster: Cluster1  
vcenter\_vmtemplate: rhel76  
vcenter\_folder\_full: /Cluster1/OCP1/

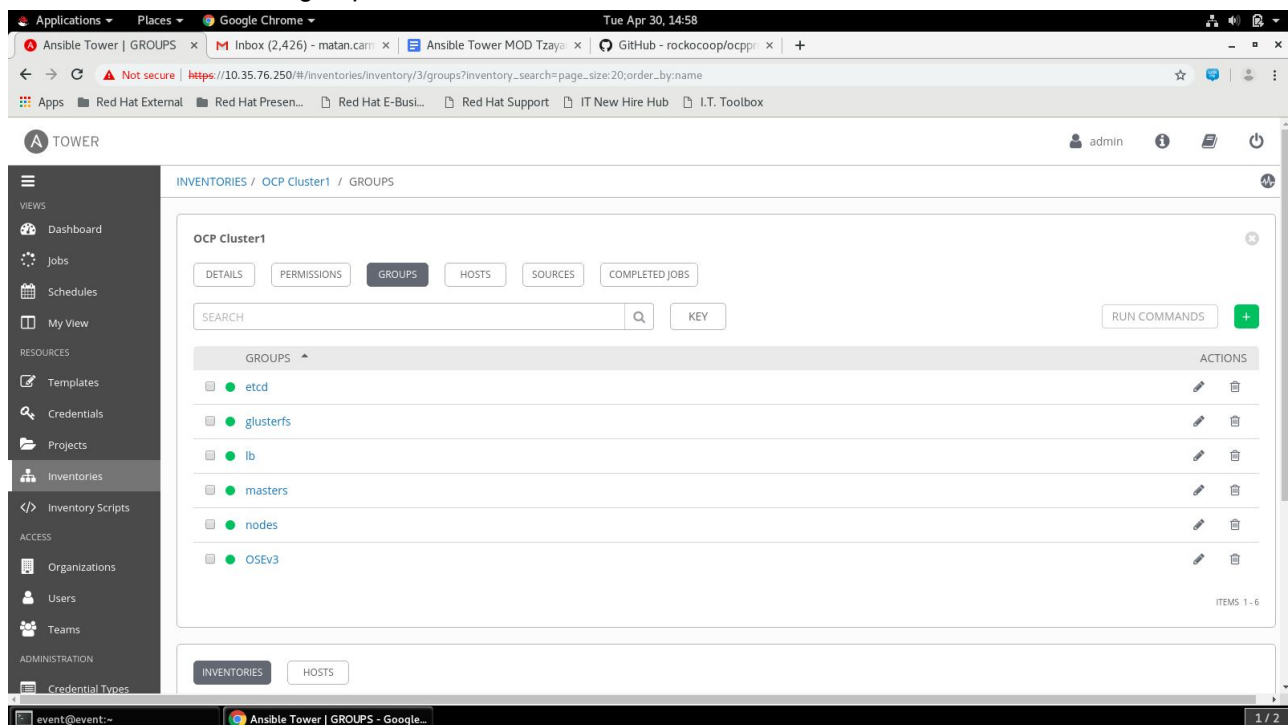
##### General Parameters #####

ansible\_ssh\_common\_args: '-o StrictHostKeyChecking=no -o ControlMaster=auto -o ControlPersist=60s -o PreferredAuthentications=publickey'  
ansible\_become: false  
cluster\_timezone: UTC  
#ntp=["10.56.190.1", "10.56.190.2"]  
searchdomain: ocp1.zayadtest.com  
dns: ["10.35.76.249", "8.8.8.8"]  
openshift\_docker\_insecure\_registries: 10.35.76.249:5000  
openshift\_docker\_blocked\_registries: registry.access.redhat.com,docker.io  
yumrepo\_url: "10.35.76.249/ocprepo"  
keepalived\_vip: 10.35.76.209  
keepalived\_interface: ens192  
keepalived\_vrrpid: 1  
routervialb: true  
Roottemppas: Password@123  
  
#connect to external grafana  
grafanaURL: "http://10.35.76.246:3000"  
grafanaPass: "Password@123"  
grafanaClusterName: "OCP1"  
grafmas: "ocpmas1.ocp1.zayadtest.com"  
prometheusURL: "https://prometheus-k8s-openshift-monitoring.apps.openshift.ocp1.zayadtest.com"

And it looked like this:



After this we created the groups of the hosts:





- Etcd: this group contains the hosts that was meant to serve as etcd, the parameter that relates to the etcd hosts:  
---  
pv\_etcd: sdc
- Glusterfs: this group contains the hosts that was meant to serve as glusterfs, it doesn't have any special variables that relates to the glusterfs hosts.
- Lb: this group contains the hosts that was meant to serve as load balancer, the parameter that relates to the lb hosts:  
---  
vmCPUs: 2  
vmMemory: 4096  
vmDisk: 40
- Masters: this group contains the hosts that was meant to serve as masters, it doesn't have any special variables that relates to the masters hosts.
- Nodes: this group contains all the hosts in the cluster except of the load balancers, the parameter that relates to the nodes hosts:  
---  
vmCPUs: 4  
vmMemory: 16384  
vmDisk: 40  
pv\_device: sdb
- OSEv3: this group contains the all the other groups (if you configures that the etcd role will be on the masters so you don't need to put the etcd group in the OSEv3 group ), the parameter that relates to the OSEv3 hosts:  
---  
#####General Parameters#####  
ansible\_ssh\_user: root  
ansible\_become: true  
debug\_level: 2  
openshift\_deployment\_type: openshift-enterprise  
openshift\_image\_tag: v3.11.98  
openshift\_release: 3.11.98  
openshift\_pkg\_version: -3.11.98  
openshift\_disable\_check: memory\_availability,disk\_availability,docker\_image\_availability  
openshift\_master\_default\_subdomain: apps.openshift.ocp1.zayadtest.com  
openshift\_master\_cluster\_hostname: openshift.ocp1.zayadtest.com  
openshift\_master\_cluster\_public\_hostname: openshift.ocp1.zayadtest.com  
  
#####The identity providers using HTTPasswd#####  
openshift\_master\_identity\_providers: [{'name': 'httpasswd\_auth', 'login': 'true', 'challenge': 'true', 'kind': 'HTPasswdPasswordIdentityProvider'}]  
openshift\_master\_htpasswd\_users: {'admin': '\$apr1\$udcx1.9E\$XZmOZkWhIRa5lzxlcBOxz1'}
- #####Disconnected environment#####  
oreg\_url: 10.35.76.249:5000/openshift3/ose-\${component}:\${version}

```
openshift_examples_modify_imagestreams: True
openshift_enable_service_catalog: False
template_service_broker_install: True
#openshift_template_service_broker_namespaces: ['openshift']
#ansible_service_broker_install: true
#ansible_service_broker_local_registry_whitelist: ['.*-apb$']
os_firewall_use_firewalld: True
openshift_hosted_manage_registry: True
openshift_cluster_monitoring_operator_install: True
openshift_metrics_install_metrics: True
openshift_logging_install_logging: True
openshift_logging_es_nodeselector: {"node-role.kubernetes.io/infra": "true"}
openshift_logging_es_memory_limit: 2Gi
openshift_node_groups: [{ 'name': 'node-config-master', 'labels':
['node-role.kubernetes.io/master=true', 'runtime=docker'] }, { 'name': 'node-config-infra', 'labels':
['node-role.kubernetes.io/infra=true', 'runtime=docker', 'logging-infra=elasticsearch'] }, { 'name': 'node-co
nfig-compute', 'labels': ['node-role.kubernetes.io/compute=true', 'runtime=docker'], 'edits': [{
'key': 'kubeletArguments.pods-per-core', 'value': ['20']}]}]
Network sdn plugin changed to openshift-ovs-multitenant
os_sdn_network_plugin_name: redhat/openshift-ovs-multitenant

Vsphere backend pv
openshift_cloudprovider_kind: vsphere
openshift_cloudprovider_vsphere_username: administrator@vsphere.local
openshift_cloudprovider_vsphere_password: Password@123
openshift_cloudprovider_vsphere_host: 10.35.76.242
openshift_cloudprovider_vsphere_datacenter: Datacenter1
openshift_cloudprovider_vsphere_cluster: Cluster1
openshift_cloudprovider_vsphere_datastore: datastore3
openshift_cloudprovider_vsphere_folder: /Cluster1/OCP1/

Set the metrics dynamic storage
openshift_master_dynamic_provisioning_enabled: True
openshift_metrics_storage_kind: dynamic
openshift_metrics_cassandra_storage_type: dynamic
openshift_metrics_cassandra_pvc_size: 15Gi
Set the logging dynamic storage
openshift_logging_storage_kind: dynamic
openshift_logging_es_pvc_dynamic: True
openshift_logging_es_pvc_size: 20Gi
openshift_logging_storage_access_modes: ['ReadWriteOnce']
Set the registry dynamic storage
openshift_hosted_registry_storage_kind: dynamic
openshift_hosted_registry_storage_access_modes: ['ReadWriteOnce']
```

```
openshift_hosted_registry_storage_volume_size: 10Gi
Set the prometheus dynamic storage
openshift_prometheus_storage_kind: dynamic
openshift_cluster_monitoring_operator_prometheus_storage_enabled: true
openshift_cluster_monitoring_operator_prometheus_storage_capacity: 50Gi
Set the prometheus-alertmanager dynamic storage
openshift_prometheus_alertmanager_storage_kind: dynamic
openshift_cluster_monitoring_operator_alertmanager_storage_enabled: true
openshift_cluster_monitoring_operator_alertmanager_storage_capacity: 2Gi
Set the prometheus-alertbuffer dynamic storage
openshift_prometheus_alertbuffer_storage_kind: dynamic
openshift_prometheus_alertbuffer_storage_type: pvc
openshift_prometheus_alertbuffer_pvc_size: 10Gi
openshift_prometheus_pvc_access_modes: [ReadWriteOnce]

if you want to use also with glusterfs as a pv you need to uncomment it.
#openshift_storage_glusterfs_image: 10.35.76.249:5000/rhgs3/rhgs-server-rhel7:v3.11
#openshift_storage_glusterfs_block_image:10.35.76.249:5000/rhgs3/rhgs-gluster-block-prov-rhel7:v3.11
#openshift_storage_glusterfs_heketi_image: 10.35.76.249:5000/rhgs3/rhgs-volmanager-rhel7:v3.11
```

Initially we create the templates that defines which playbook will run. We have several templates:

- Deploy VMware: this template runs the playbooks/provisioningVMware.yml from vmware-ansible project, this playbook run 3 roles:
  1. VMware\_create\_vm: this role creates the cluster vms, the main.yml looks like this:

---

```
- name: Abort Play when IP is in use
 command: ping -c1 {{ ansible_ssh_host }}
 delegate_to: localhost
 register: ping_result
 failed_when: >
 (" " not in ping_result.stderr)
 changed_when: False

- meta: end_play
 when: ping_result.rc==0

- name:
 debug:
 msg: " IP Address {{ ansible_ssh_host }} is free"
 verbosity: 3
 when: ping_result.rc==1
```

```
- name: Create a new VM
block:
 - set_fact:
 disk_size: "{{ vmDisk }}"
 memory: "{{ vmMemory }}"
 cpu: "{{ vmCPUs }}"

 - vmware_guest:
 hostname: "{{ vcenter_hostname }}"
 username: "{{ vcenter_user }}"
 password: "{{ vcenter_pass }}"
 datacenter: "{{ vcenter_datacenter }}"
 name: "{{ hostname }}"
 template: "{{ vcenter_vmtemplate }}"
 resource_pool: "{{ vcenter_resource_pool | default(omit) }}"
 cluster: "{{ vcenter_cluster }}"
 state: poweredon
 folder: "{{ vcenter_folder_full }}"
 validate_certs: False
 is_template: no
 disk:
 - size_gb: "{{ disk_size }}"
 type: thin
 datastore: "{{ vcenter_datastore }}"
 hardware:
 memory_mb: "{{ memory }}"
 num_cpus: "{{ cpu }}"
 networks:
 - name: "{{ vlan }}"
 ip: "{{ ansible_ssh_host }}"
 netmask: "{{ netmask }}"
 gateway: "{{ gateway }}"
 customization:
 dns_servers:
 - "{{ dns[0] }}"
 - "{{ dns[1] }}"
 wait_for_ip_address: yes
 register: new_vm
 delegate_to: localhost
 failed_when: >
 (new_vm.failed==true)
 rescue:
 - debug:
 msg: " VM {{ hostname }} already Exists"
 verbosity: 3
```

```
- meta: end_play
```

```
- name: New VM Status Message
```

```
 debug:
```

```
 msg: " New VM {{ hostname }}, {{ ansible_ssh_host }}, is Up and Running"
```

2. Vmware\_add\_disk: this role adds a disk to all the cluster nodes except the lb hosts, this is the main.yml:

```

```

```
- name: Load govc Vars
```

```
 include_vars: "{{ role_path }}/vars/govc_vars.yml"
```

```
- name: Get VM Info/State
```

```
 command: "govc vm.info -vm.ip={{ ansible_ssh_host }} -k=true"
```

```
 environment:
```

```
 GOVC_USERNAME: "{{ govc_username }}"
```

```
 GOVC_PASSWORD: "{{ govc_password }}"
```

```
 GOVC_INSECURE: "{{ govc_insecure }}"
```

```
 GOVC_DATASTORE: "{{ govc_datastore }}"
```

```
 GOVC_DATACENTER: "{{ govc_datacenter }}"
```

```
 GOVC_URL: "{{ govc_url }}"
```

```
 register: govc_vm_info
```

```
 changed_when: False
```

```
 delegate_to: localhost
```

```
 failed_when: >
```

```
 ("no such VM" in govc_vm_info.stdout)
```

```
- name: Add SCSI to existing VM
```

```
 command: "govc device.scsi.add -vm={{ hostname }} -type=pvscsi
```

```
-k=true"
```

```
 environment:
```

```
 GOVC_USERNAME: "{{ govc_username }}"
```

```
 GOVC_PASSWORD: "{{ govc_password }}"
```

```
 GOVC_INSECURE: "{{ govc_insecure }}"
```

```
 GOVC_DATASTORE: "{{ govc_datastore }}"
```

```
 GOVC_DATACENTER: "{{ govc_datacenter }}"
```

```
 GOVC_URL: "{{ govc_url }}"
```

```
 register: govc_scsi_add
```

```
 changed_when: False
```

```
 delegate_to: localhost
```

```
 failed_when: >
```

```
 ("pvscsi-100" not in govc_scsi_add.stdout)
```

```
 with_items:
```

```
 - "{{ disks }}"
```

```
 loop_control:
```

```

 loop_var: scsi

- name: wait for reconfiguration
 pause: seconds=10

- name: Add Disk to VM
 block:
 # Add /dev/sdb Disk to existing VM
 - command: >
 govc vm.disk.create
 -vm={{ hostname }}
 -name='{{ hostname }}/{{ hostname }}_{{index}}'
 -ds='{{ vcenter_datacenter }}/datastore/{{ vcenter_datastore
 }}'

 -mode=independent_persistent
 -controller pvscsi-100{{index + 1 }}
 -eager=false
 -thick=false
 -size {{ item }}G
 -k=true
 environment:
 GOVC_USERNAME: "{{ govc_username }}"
 GOVC_PASSWORD: "{{ govc_password }}"
 GOVC_INSECURE: "{{ govc_insecure }}"
 GOVC_DATASTORE: "{{ govc_datastore }}"
 GOVC_DATACENTER: "{{ govc_datacenter }}"
 GOVC_URL: "{{ govc_url }}"
 register: govc_sdb_add
 changed_when: False
 delegate_to: localhost
 failed_when: >
 ("Creating disk" not in govc_sdb_add.stdout)
 with_items:
 - "{{disks}}"
 loop_control:
 index_var: index
3. Vm_power: this role restart all the vms, the main.yml:

- name: Power Control
 vmware_guest:
 hostname: "{{ vcenter_hostname }}"
 username: "{{ vcenter_user }}"
 password: "{{ vcenter_pass }}"
 validate_certs: no
 cluster: "{{ vcenter_cluster }}"
 name: "{{ hostname }}"
 state: "{{ vm_state }}"

```

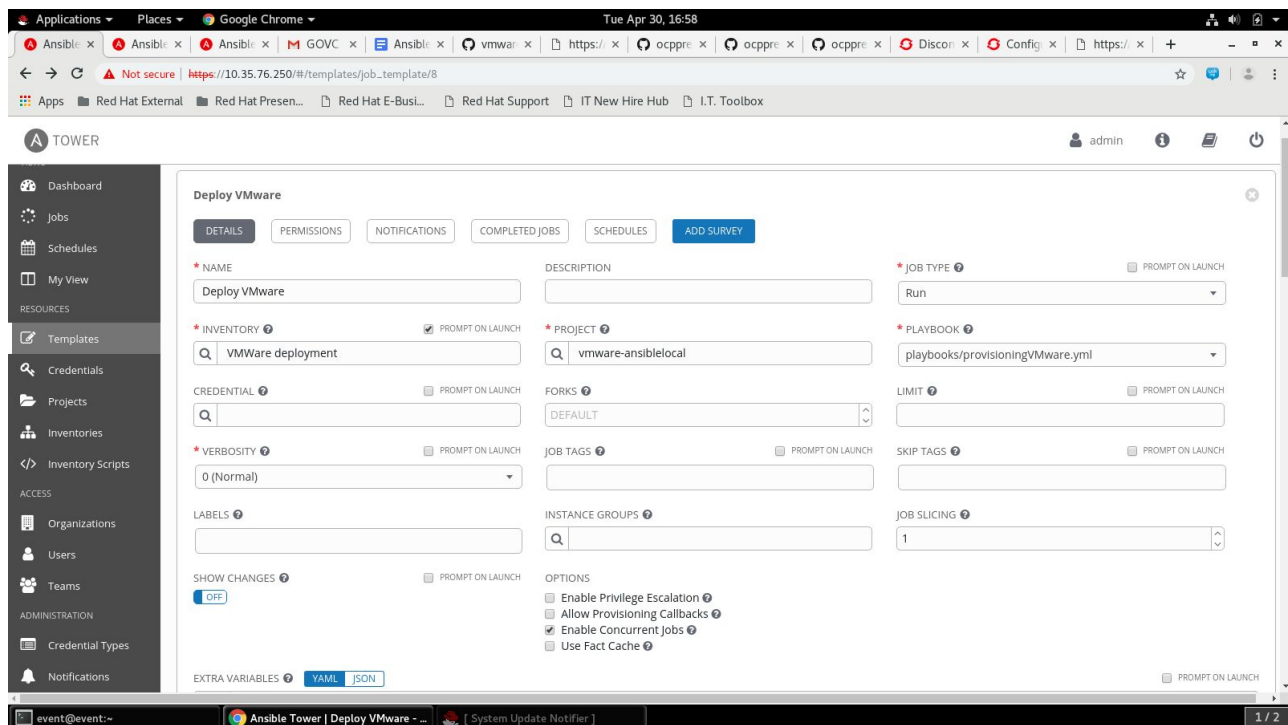
```

delegate_to: localhost

- pause: seconds=30

#- name: Wait 300 seconds for remote system to be available, but only
start checking after 30 seconds
wait_for_connection:
delay: 30
timeout: 300

```



- The second template is Destroy VMware, this template runs the playbook `playbooks/destroyEnvironment.yml` from `vmware-ansible` project, this playbook runs 3 roles:
  1. `Vm_power`: this role powers off all the cluster VMs, this is the `main.yml`:

---

```

- name: Power Control
 vmware_guest:
 hostname: "{{ vcenter_hostname }}"
 username: "{{ vcenter_user }}"
 password: "{{ vcenter_pass }}"
 validate_certs: no
 cluster: "{{ vcenter_cluster }}"
 name: "{{ hostname }}"
 state: "{{ vm_state }}"
 delegate_to: localhost

```

- pause: seconds=30

#- name: Wait 300 seconds for remote system to be available, but only start checking after 30 seconds

# wait\_for\_connection:

# delay: 30

# timeout: 300

2. Vmware\_remove\_vm: this role erase all the cluster vms, this is the main.yml:

- name: Load govc Vars

include\_vars: "{{ role\_path }}/vars/govc\_vars.yml"

- name: Remove "{{ hostname }}"

vmware\_guest:

hostname: "{{ vcenter\_hostname }}"

username: "{{ vcenter\_user }}"

password: "{{ vcenter\_pass }}"

validate\_certs: no


cluster: "{{ vcenter\_cluster }}"

name: "{{ hostname }}"

state: absent

delegate\_to: localhost

register: facts



The screenshot shows the Ansible Tower web interface in Google Chrome. The browser address bar shows a URL for a job template search. The interface has a sidebar with navigation options like Dashboard, Jobs, Schedules, My View, Templates, Credentials, Projects, Inventories, Inventory Scripts, Organizations, Users, Teams, and Administration. The main content area displays the configuration for a job named 'Destory VMware'. The configuration includes fields for Name, Description, Job Type (set to 'Run'), Inventory (set to 'VMWare deployment'), Project (set to 'vmware-ansiblelocal'), Playbook (set to 'playbooks/destroyEnvironmet.yml'), Credential, Forks (set to 'DEFAULT'), Limit, Verbosity (set to '0 (Normal)'), Job Tags, Skip Tags, Labels, Instance Groups, Job Slicing (set to '1'), Show Changes (set to 'OFF'), and Options (with checkboxes for 'Enable Privilege Escalation', 'Allow Provisioning Callbacks', 'Enable Concurrent Jobs', and 'Use Fact Cache'). There are also tabs for DETAILS, PERMISSIONS, NOTIFICATIONS, COMPLETED JOBS, SCHEDULES, and ADD SURVEY.

- The third template is Pre-Install, this template run the playbook playbooks/pre-install.yml from ocpprep project, this playbook performs pre install tasks on the cluster nodes, this playbooks runs this playbooks:



1. Setup\_internal\_repos.yml: this playbook configure all the repository sources for openshift in the cluster servers, the main.yml:

---

- name: mkdir oldrepos  
file:  
  path: /etc/yum.repos.d/.oldrepos  
  state: directory  
  tags: mkdir\_oldrepos
- name: Check to see if any repo files found in /etc/yum.repos.d  
find:  
  paths: "/etc/yum.repos.d"  
  patterns: "\*.repo"  
  register: repo\_files
- name: Move all yum.repo files to /etc/yum.repos.d/.oldrepos  
shell: "mv /etc/yum.repos.d/\*.repo /etc/yum.repos.d/.oldrepos/"  
tags: mv\_repo  
when: repo\_files.matched|int > 0
- name: RHEL Openshift repos setup  
template:  
  src: templates/setup\_internal\_repos.yml  
  dest: /etc/yum.repos.d/ose.repo

2. Then it installs the supported packages on all the servers:

- name: install support packages  
package:  
  name: "{{item}}"  
  state: present  
with\_items:  
  - wget  
  - git  
  - net-tools  
  - bind-utils  
  - yum-utils  
  - iptables-services  
  - bridge-utils  
  - bash-completion  
  - kexec-tools  
  - sos  
  - psacct  
  - atomic-openshift-docker-excluder  
  - atomic-openshift-excluder  
  - ntp


3. Then it's configuring the ntp, update resolv file for dns and builds hosts file:
- name: Configure NTP | Add server lines to ntp.conf
    - template:
      - src: templates/ntpconf.j2
      - dest: /etc/ntp.conf
      - when: ntp is defined
      - notify: ntp\_restart
  - name: Configure NTP | start ntp daemon
    - service:
      - name: ntpd
      - enabled: yes
      - state: started
  - name: update resolv file for dns
    - template:
      - src: templates/resolv.j2
      - dest: /etc/resolv.conf
      - when: dns is defined
  - name: set hostname
    - hostname:
      - name: "{{ hostname }}"
  - name: Build hosts file
    - lineinfile:
      - path: /etc/hosts
      - line: "{{ hostvars[item].ansible\_default\_ipv4.address }} {{ hostvars[item].hostname }}"
      - state: present
      - when: 'hostvars[item].ansible\_default\_ipv4.address is defined'
      - with\_items:
        - "{{ groups['OSEv3'] }}"
    - handlers:
      - name: ntp\_restart
        - service:
          - name: ntpd
          - state: restarted
    - name: update /etc/hosts on all nodes with ivpcoe-vip
      - lineinfile:
        - dest: /etc/hosts
        - line: '{{ keepalived\_vip }} {{ openshift\_master\_cluster\_public\_hostname }}'
        - when: keepalived\_vip is defined

4. Configure the keepalive on the load balancers:
  - hosts: lb
    - become: true
    - tasks:
      - name: install keepalived on lb nodes
        - package:
          - name: "{{item}}"
        - with\_items:
          - keepalived
        - when: keepalived\_vip is defined
      - name: create keepalived.conf
        - template:
          - src: ../templates/keepalived.conf
          - dest: /etc/keepalived/keepalived.conf
        - when: keepalived\_vip is defined
      - name: restart keepalived
        - service:
          - name: keepalived
          - state: restarted
          - enabled: yes
        - when: keepalived\_vip is defined
5. Initially it installs docker on all servers except the lbs:
  - name: Install Docker
    - hosts: nodes
    - become: true
    - gather\_facts: true
    - serial: "100%"
    - tasks:
      - name: Install Docker | install packages
        - package:
          - name: "{{ item }}"
          - state: present
        - with\_items:
          - docker
          - ethtool
        - notify: docker\_restart

- name: Install Docker | setup docker storage  
template:
  - src: templates/docker-storage-setup.j2
  - dest: /etc/sysconfig/docker-storage-setupnotify: docker\_restart
- name: Verify presence of /dev/docker-vg/docker-pool  
stat:
  - path: /dev/docker-vg/docker-poolregister: docker\_vg\_status
- name: Run docker-storage-setup  
command: /usr/bin/docker-storage-setup  
when: docker\_vg\_status.stat.islnk is not defined
- name: Install Docker | enable docker  
service:
  - name: docker
  - state: started
  - enabled: yes
- name: setup etcd storage  
include\_tasks: include/etcdlvol.yml  
when: inventory\_hostname in groups['etcd'] and pv\_etcd is defined

handlers:

- name: docker\_restart  
service:
  - name: docker
  - state: restarted



- The fourth template is OCP Prerequisites, it runs the playbook playbooks/Prerequisites.yml from the project openshift-ansible, this playbook runs the Prerequisites tasks that needs to be done before the deployment. This is the playbook:

---

# I\_scale\_up\_hosts may be passed in via various scaleup plays.

- name: Fail openshift\_kubelet\_name\_override for new hosts

hosts: "{{ I\_scale\_up\_hosts | default('nodes') }}"

tasks:

- name: Fail when openshift\_kubelet\_name\_override is defined

fail:

msg: "openshift\_kubelet\_name\_override Cannot be defined for new hosts"

when:

- openshift\_kubelet\_name\_override is defined

- openshift\_cloudprovider\_kind | default("", true) != 'azure'

- import\_playbook: init/main.yml

vars:

I\_install\_base\_packages: True

I\_repo\_hosts: "{{ I\_scale\_up\_hosts | default('oo\_all\_hosts') }}"

- import\_playbook: init/validate\_hostnames.yml

when: not (skip\_validate\_hostnames | default(False))

# This is required for container runtime for cri-o, only needs to run once.

- name: Configure os\_firewall


```

hosts: "{{ l_scale_up_hosts | default(l_default_firewall_hosts) }}"
vars:
 l_default_firewall_hosts:
"oo_masters_to_config:oo_etcd_to_config:oo_lb_to_config:oo_nfs_to_config:oo_nodes_to_config"
roles:
 - role: os_firewall

```

- import\_playbook: container-runtime/private/setup\_storage.yml

- import\_playbook: container-runtime/private/config.yml



- The fifth template is OCP Deploy Cluster, it runs the playbook playbooks/deploy\_cluster.yml from the project openshift-ansible, this playbook runs the deployment tasks of the openshift cluster. This is the playbook:

---

```

- import_playbook: init/main.yml
vars:
 l_networkman_check_hosts: "oo_nodes_to_config"

- import_playbook: openshift-checks/private/install.yml


- import_playbook: openshift-node/private/bootstrap.yml

- import_playbook: common/private/control_plane.yml

- import_playbook: openshift-node/private/join.yml

```

- import\_playbook: common/private/components.yml



The screenshot shows the Ansible Tower web interface for configuring a template named 'OCP Deploy Cluster'. The interface includes a sidebar with navigation options like Dashboard, Jobs, Schedules, My View, Templates, Credentials, Projects, Inventories, Inventory Scripts, Organizations, Users, Teams, and Administration. The main content area displays the configuration for the 'OCP Deploy Cluster' template, with tabs for DETAILS, PERMISSIONS, NOTIFICATIONS, COMPLETED JOBS, and SCHEDULES. The configuration fields include:

- NAME:** OCP Deploy Cluster
- DESCRIPTION:** (empty)
- JOB TYPE:** Run
- INVENTORY:** OCP Cluster1
- PROJECT:** openshift-ansible
- PLAYBOOK:** playbooks/deploy\_cluster.yml
- CREDENTIAL:** templateKey
- FORKS:** DEFAULT
- LIMIT:** (empty)
- VERBOSITY:** 0 (Normal)
- JOB TAGS:** (empty)
- SKIP TAGS:** (empty)
- LABELS:** (empty)
- INSTANCE GROUPS:** (empty)
- JOB SLICING:** 1
- SHOW CHANGES:** OFF
- OPTIONS:**
  - Enable Privilege Escalation
  - Allow Provisioning Callbacks
  - Enable Concurrent Jobs
  - Use Fact Cache

- The sixth template is Post Jobs, this template runs the playbook playbooks/post-jobs.yml from the ocpprep project, this playbook run some tasks:

1. First it creates a datasource for the cluster in the central grafana:

```

- hosts: localhost
 become: true
 gather_facts: false
 serial: 1

```

tasks:

- name: copy jq binary to master

```

copy:
 src: "../supportBinary/jq-linux64"
 dest: "/opt/jq"
 mode: 777
 delegate_to: "{{grafmas}}"
 when: grafmas is defined

```

- name: retrieve password

```

shell: oc get secret grafana-datasources -n openshift-monitoring -o json | /opt/jq
'.data["prometheus.yaml"]' | tr -d '"' | base64 -d | grep -i basicAuthPassword | awk '{print $2}' | tr
-d ' ' | tr -d '"'
register: promPass

```

```
delegate_to: "{{grafmas}}"
when: grafmas is defined
```

```
- name: Create datasource for cluster
 grafana_datasource:
 name: "{{grafanaClusterName}}"
 grafana_url: "{{grafanaURL}}"
 grafana_user: "admin"
 grafana_password: "{{grafanaPass}}"
 ds_type: "prometheus"
 url: "{{prometheusURL}}"
 basic_auth_user: "internal"
 basic_auth_password: "{{promPass.stdout}}"
 validate_certs: False
 tls_skip_verify: true
 when: grafmas is defined
```

2. Then it adds some configuration for the load balancers:

```
- hosts: lb
 become: true
 gather_facts: false
 serial: 1
```

tasks:


```
- name: open port 80
 firewallld:
 port: 80/tcp
 permanent: true
 immediate: yes
 state: enabled

- name: open port 443
 firewallld:
 port: 443/tcp
 permanent: true
 immediate: yes
 state: enabled

- name: open port 2379
 firewallld:
 port: 2379/tcp
 permanent: true
 immediate: yes
 state: enabled
```




- name: restart firewalld  
service:
  - name: firewalld
  - state: restarted
  
- name: update haproxy config on lb  
template:
  - src: templates/haproxy.cfg
  - dest: haproxy.cfg
  - when: routervialb is defined and routervialb == true
  - delegate\_to: localhost
  
- name: copy haproxy cfg to lb  
copy:
  - src: haproxy.cfg
  - dest: /etc/haproxy/
  
- name: selinux to open 2379  
seport:
  - ports: 2379
  - proto: tcp
  - setype: http\_port\_t
  - state: present
  
- name: restart haproxy  
service:
  - name: haproxy
  - state: restarted



- The seventh template is OCP Deploy service catalog, this template runs the playbook `playbooks/openshift-service-catalog/config.yml` from the `openshift-ansible` project. This playbook runs the deployment of the service catalog:


---

- `import_playbook: ../init/main.yml`  
`vars:`
  - `l_init_fact_hosts: "oo_masters_to_config:oo_etcd_to_config"`
  - `l_openshift_version_set_hosts: "oo_masters_to_config:!oo_first_master"`
  - `l_sanity_check_hosts: "{{ groups['oo_masters_to_config'] }}"`
- `import_playbook: private/config.yml`



The screenshot shows the Ansible Tower web interface. The left sidebar contains navigation links for Views (Dashboard, Jobs, Schedules, My View), Resources (Templates, Credentials, Projects, Inventories, Inventory Scripts), Access (Organizations, Users, Teams), and Administration (Credential Types). The main content area is titled 'TEMPLATES / OCP Deploy service catalog'. It features a 'DETAILS' tab and buttons for 'PERMISSIONS', 'NOTIFICATIONS', 'COMPLETED JOBS', 'SCHEDULES', and 'ADD SURVEY'. The configuration form includes fields for NAME ('OCP Deploy service catalog'), DESCRIPTION, JOB TYPE ('Run'), INVENTORY ('OCP Cluster1'), PROJECT ('openshift-ansible'), PLAYBOOK ('playbooks/openshift-service-catalog/config.yml'), CREDENTIAL ('templateKey'), FORKS ('DEFAULT'), LIMIT, VERBOSITY ('0 (Normal)'), JOB TAGS, SKIP TAGS, JOB SLICING ('1'), and SHOW CHANGES ('OFF'). There are also checkboxes for 'PROMPT ON LAUNCH' and 'ENABLE CONCURRENT JOBS', and a section for 'OPTIONS' including 'Enable Privilege Escalation', 'Allow Provisioning Callbacks', 'Enable Concurrent Jobs', and 'Use Fact Cache'.

- The last template is VMware and OCP Full Deployment, this is a template from kind workflow, that means that it runs other jobs in order that you choose:

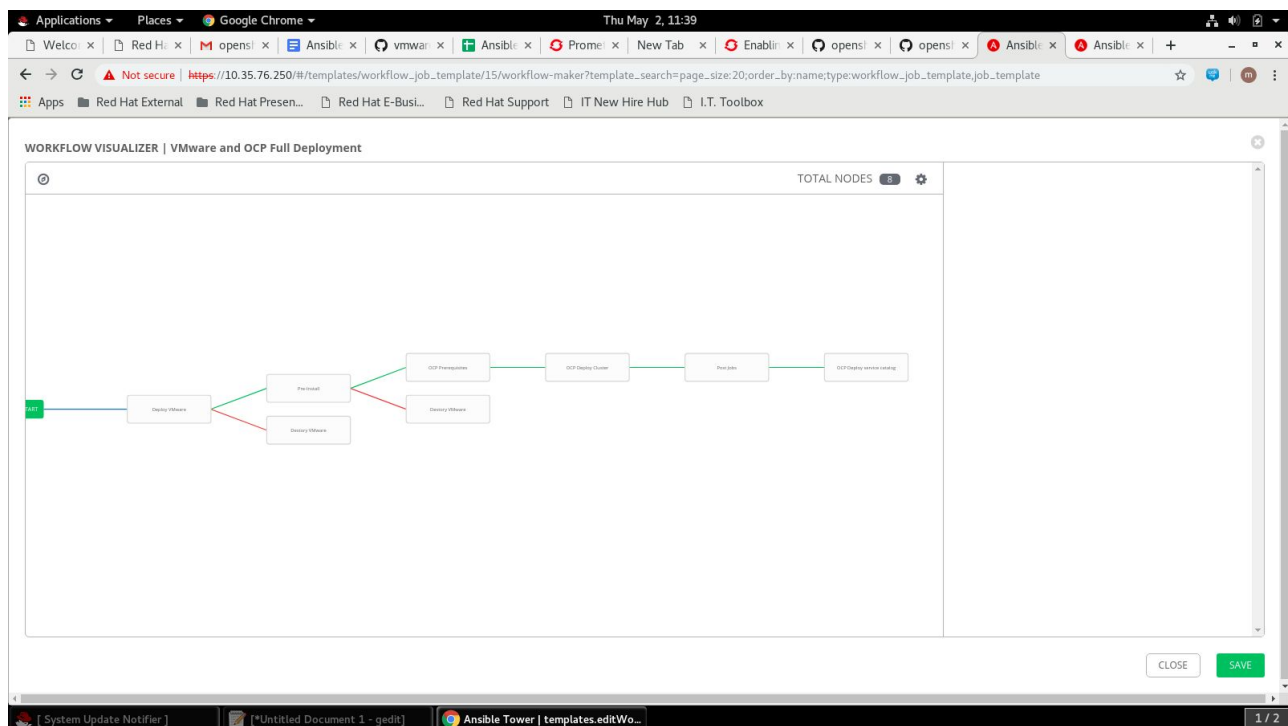


The screenshot shows the Ansible Tower web interface for the 'VMware and OCP Full Deployment' template. The left sidebar is the same as the previous screenshot. The main content area is titled 'TEMPLATES / VMware and OCP Full Deployment'. It features a 'DETAILS' tab and buttons for 'PERMISSIONS', 'NOTIFICATIONS', 'COMPLETED JOBS', 'SCHEDULES', 'ADD SURVEY', and 'WORKFLOW VISUALIZER'. The configuration form includes fields for NAME ('VMware and OCP Full Deployment'), DESCRIPTION, ORGANIZATION, INVENTORY ('OCP Cluster1'), LABELS, and EXTRA VARIABLES (with tabs for 'YAML' and 'JSON'). There is a checkbox for 'ENABLE CONCURRENT JOBS' and a 'SHOW CHANGES' section. At the bottom, there are 'CANCEL' and 'SAVE' buttons, and a 'TEMPLATES' section with a search bar and a '+ KEY' button.

The order of the workflow is: first you run the Deploy VMware template, if it fails it runs the Destroy VMware template and exit. If the Deploy VMware works the workflow continues to template Pre-Install, if it fails it

runs the Destroy VMware template and exit. if the Pre-Install works the workflow continues to template OCP Prerequisites, if it fails the main template stop, if it works it continues to template OCP Deploy Cluster, if it fails the main template stop, if it works it continues to template Post Jobs, if it fails the main template stop, if it works it continues to template OCP Deploy service catalog, if it fails the main template stop, if it works that means that the workflow template worked.

And this is the workflow:



## 5. ENGAGEMENT CLOSURE AND RECOMMENDATIONS.

### 5.1 Knowledge Transfer

For this engagement the knowledge transfer was provided during the engagement. On this environment knowledge transfer was specifically provided to the following individual(s):

| Client Staff Member | Notes |
|---------------------|-------|
|---------------------|-------|

|                                                     |                                                                                                                                                                                                                                                                                                                                                             |
|-----------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| NAME_OF_MOD<br>Mamram_STAFF_MEMBER<br>Matan Carmeli | <ul style="list-style-type: none"><li>Assisted with the data gathering process for the environment.</li><li>Provided all the network details for the environment (including physical and virtual).</li><li>Served as our full team member and learned as engagement was in progress</li><li>Took active part in this journal preparation and demo</li></ul> |
|-----------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|

**Table 5-1: Knowledge Transfer.**

## 5.2 Testing

Testing was conducted in accordance with a predefined test plan as documented in Appendix 7.1.

Testing was conducted by the consultant and witnessed by a representative of MOD - Mamram as follows:

| Description                | Value              |
|----------------------------|--------------------|
| Overall Test Results       | Passed             |
| Testing completed on date: | 01/05 - 02/05-2019 |
| Testing Witnessed by:      | 01/05 - 02/05-2019 |

**Table 4-2: Knowledge Transfer.**

## 5.3 Engagement Observations

Overall, the engagement went smoothly and provided a good opportunity for the Customer's team member to get hands on experience with Red Hat technologies and support practices. These skills can be improved by participation in Red Hat training programs.

## 5.4 Training and Certification

Specific course descriptions and availability can be found on Red Hat's website under Training and Certification. The following courses are recommended:

### Red Hat Certified System Administrator (RHCSA)

Verizon should invest in certifying two team members as RHCSA's within their IT support staff in order to ensure the team can adequately manage and maintain the expanding environment. Various courses are available from Red Hat to train RHCSA candidates and prepare them for the certification exam, including ground up and fast tracks for existing Unix and Linux admins – The following course is a likely fit for Example.com's experienced UNIX admin staff:

RHCSA Rapid Track course with exam (RH200)

*"The RHCSA Rapid Track course with exam (RH200) is designed for experienced Linux and Unix system administrators who want to become accredited with the RHCSA certification. Students will learn to manage a Linux server, including installation and configuration of local components and services, as well as connections to existing network services. To successfully navigate this accelerated course, students must already have solid command line skills and know how to access man pages for help."*

## Red Hat Automation with Ansible 1 (DO 407)

### Red Hat Automation with Ansible 2: Ansible Tower (DO409)

#### Course Objectives

- *Deploy and use Red Hat Ansible Tower to manage existing Ansible projects, playbooks, and roles*
- *Use the visual dashboard to centrally launch, control, and monitor Ansible jobs*
- *Configure users and teams and use them to control access to systems, projects, and other resources through role-based access controls*
- *Automatically schedule Ansible jobs and update the host inventory*
- *Perform basic maintenance and administration of the Red Hat Ansible Tower installation*
- *Use the Red Hat Ansible Tower API to launch jobs from existing templates*

## 6. APPENDIXES

### 6.1 Appendix A: RH Ansible Tower and OCP 3.11 Accelerator Test Matrix and Test result

| ID | Test                                        | Method           | Expected Result                                                                | Notes                                                                 | Pass/<br>Fail |
|----|---------------------------------------------|------------------|--------------------------------------------------------------------------------|-----------------------------------------------------------------------|---------------|
| S1 | Ansible Tower Define Job Template           | CLI or Dashboard | Template should be created successfully                                        |                                                                       | Pass          |
| S2 | Ansible Tower - Define Project              | CLI or Dashboard | Projects should be created successfully                                        |                                                                       | Pass          |
| S3 | Ansible Tower - Inventory files editing     | CLI or Dashboard | Different Inventory files should be edited successfully                        |                                                                       | Pass          |
| S4 | Ansible Tower - Define Workflow template    | CLI or Dashboard | Template should be created successfully                                        |                                                                       | Pass          |
| S5 | Ansible Tower - Running Jobs simultaneously | CLI or Dashboard | Two workflow jobs running simultaneously against two different inventory files | When working with Source GitHub - the same job run in sequential mode | Pass          |
| S6 | VMDK as persistent storage provisioning     | CLI or Dashboard | Persistent Volume with storage class vsphere-storage                           |                                                                       | Pass          |

|     |                                                                                                        |                  |                                                                                                               |                                                                            |      |
|-----|--------------------------------------------------------------------------------------------------------|------------------|---------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------|------|
|     |                                                                                                        |                  | should be created,<br>PVC should be created                                                                   |                                                                            |      |
| S7  | GlusterFS as persistent storage provisioning                                                           | CLI or Dashboard | Persistent Volume with storage class glusterfs-storage should be created, PVC should be created               |                                                                            | Pass |
| S8  | Two OCP cluster should be created including VM creation, preparation, OCP deployment, postinstall jobs | CLI or Dashboard | Workflow Templates should be completed successfully and as a result two ocp clusters should be up and running |                                                                            | Pass |
| S9  | OCP - Admin role should be provided to admin User                                                      | CLI or Dashboard | Admin user should be able to manage cluster                                                                   |                                                                            | Pass |
| S10 | OCP - New projects creation                                                                            | CLI or Dashboard | New Projects should be created successfully                                                                   |                                                                            | Pass |
| S11 | OCP - Dynamic PVC should be provisioned                                                                | CLI or Dashboard | Dynamic PVC should be created successfully                                                                    |                                                                            | Pass |
| S12 | OCP - new app creation                                                                                 | CLI or Dashboard | At least one App should be created successfully                                                               | Image streams should be edited before with definition of insecure registry | Pass |
| S13 | Ansible Tower - both Clusters should be connected to the Central Grafana                               | CLI or Dashboard | In Central Grafana both OCP Clusters should be visualized successfully.                                       |                                                                            | Pass |

OTHER APPENDIX GOES HERE.

[Copying-a-full-git-repo](#)

[https://docs.openshift.com/container-platform/3.11/install\\_config/configuring\\_vsphere.html](https://docs.openshift.com/container-platform/3.11/install_config/configuring_vsphere.html)

<https://grafana.com/docs/installation/rpm/>

<https://www.digitalocean.com/community/tutorials/how-to-use-prometheus-to-monitor-your-centos-7-server>

<https://docs.vmware.com/en/VMware-vSphere/6.7/com.vmware.vsphere.html.hostclient.doc/GUID-ED3ECA21-5763-4919-8947-A819A17980FB.html>

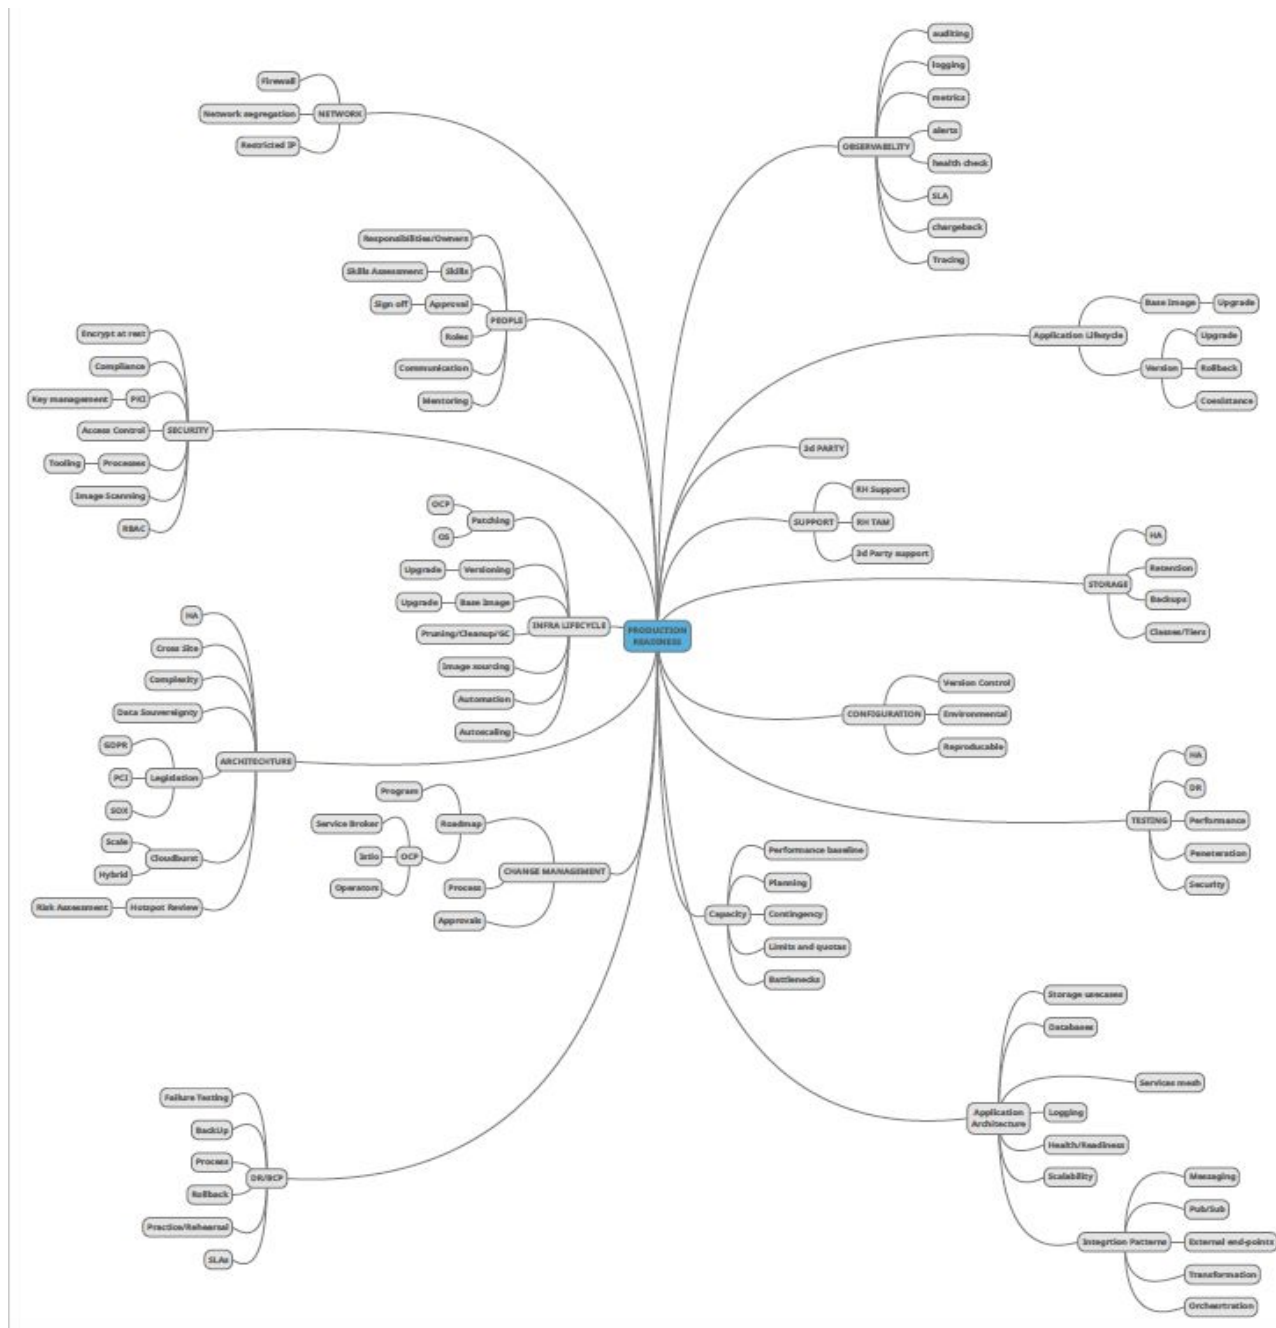
<https://access.redhat.com/articles/3344101>

[EOD Reports](#)

[Worksheet](#)



## Production Readiness Map



## PRODUCTION READINESS

**OBSERVABILITY**

Auditing

logging

metrics

alerts

health check

SLA

chargeback

Tracing

**Application Lifecycle****Base Image**

- Upgrade

**Version**

- Upgrade
- Rollback
- Co-existence

**3d PARTY**

Add all relevant third party tools and components

**STORAGE**

HA

Retention

Backups

Classes/Tiers

**SUPPORT**

RH Support

RH TAM

3d Party support

**CONFIGURATION****Version Control**

Configuration as a code management and versioning

**Environmental**

Environment specific configuration changed across application stages - DEV, Release, UAT, etc.

**Reproducible**

Support to reproduce or recreate environment

## TESTING

HA

DR

Performance

Penetration

Security

## Capacity

Performance baseline

Planning

Contingency

Limits and quotas

Bottlenecks

## Application Architecture

Storage use cases

Databases

Services mesh

Logging

Health/Readiness

Scalability

Integration Patterns

- Messaging
- Pub/Sub
- External end-points
- Transformation
- Orchestration

## SECURITY

Encrypt at rest

Compliance

PKI

- Key management

Access Control

Processes

- Tooling

Image Scanning

RBAC

## NETWORK

Firewall

Network segregation

Restricted IP

## INFRA LIFECYCLE

Patching

- OCP
- OS

Versioning

- Upgrade

Base Image

- Upgrade

Pruning/Cleanup/GC

Image sourcing

Automation

Autoscaling

## PEOPLE

Responsibilities/Owners

Skills

Skills Assessment

Approval

Sign off

Roles

Communication

Mentoring

## ARCHITECTURE

HA

Cross Site

Complexity

Data Sovereignty

Legislation

- GDPR
- PCI
- SOX

Cloudburst

- Scale
- Hybrid

Hotspot Review

- Risk Assessment

## **CHANGE MANAGEMENT**

### **Roadmap**

- Program
- OCP
- Service Broker
- Istio
- Operators

### **Process**

### **Approvals**

## **DR/BCP**

### **Failure Testing**

### **BackUp**

### **Process**

### **Rollback**

### **Practice/Rehearsal**

### **SLAs**